

# Summary of My Research

- **Name:** Jianqing Zhang
- **Ph.D.:** Shanghai Jiao Tong & Tsinghua University
- **Visiting:** Hong Kong Polytechnic University
- **Home Page:** [github.com/TsingZ0](https://github.com/TsingZ0)
- **E-mail:** [tsingz@sjtu.edu.cn](mailto:tsingz@sjtu.edu.cn)
- **LinkedIn:** [www.linkedin.com/in/tsingz/](https://www.linkedin.com/in/tsingz/)
- **X:** @TsingZ00





# Overview

- **Research interests:** Domain Adaptation (**core**), Model Merging, Personalization
- **Fields (transfer ability):** *Code LLM, Synthetic Data Generation, Distributed Learning, Recommender System*
- **Outstanding advantages:** Research capabilities and engineering experience
- **Publications:** 9 first-author top-tier conference/journal papers
  - Stage ① [AAAI'23 \(oral\)](#), [KDD'23](#), [ICCV'23](#), [NeurIPS'23](#), [JMLR'25](#)
  - Stage ② [AAAI'24](#), [CVPR'24](#), [KDD'25](#)
  - Stage ③ [EMNLP'24](#), [ICML'25](#), [ICML'25 \(spotlight\)](#)
- **Open-sourced projects (initiator):**
  - [EvolveGen](#), [PFLlib](#) (**1800+** stars, **300+** forks), [HtFLlib](#), HtFLlib on Device, [FL-IoT](#), etc.
- **Awards:** Youth Talent of China Association for Science and Technology (Chinese Association for Artificial Intelligence, CAAI), Wenjun Wu Honorary Doctorate in AI, PhD National Scholarship
- **Projects:** ① Cross-hospital cancer recognition, ② Cross-province intelligent 12345 hotline model, ③ HtFL testbed on real-world devices, ④ Led 9-member team in building a distributed ML platform
- **Impact:** **700+** citations, **30K+** views across major media, well-recognized by IEEE/ACM Fellows
- **Intern:** ByteDance AML, Tsinghua AIR, KAUST SANDs lab, Tencent AI Code



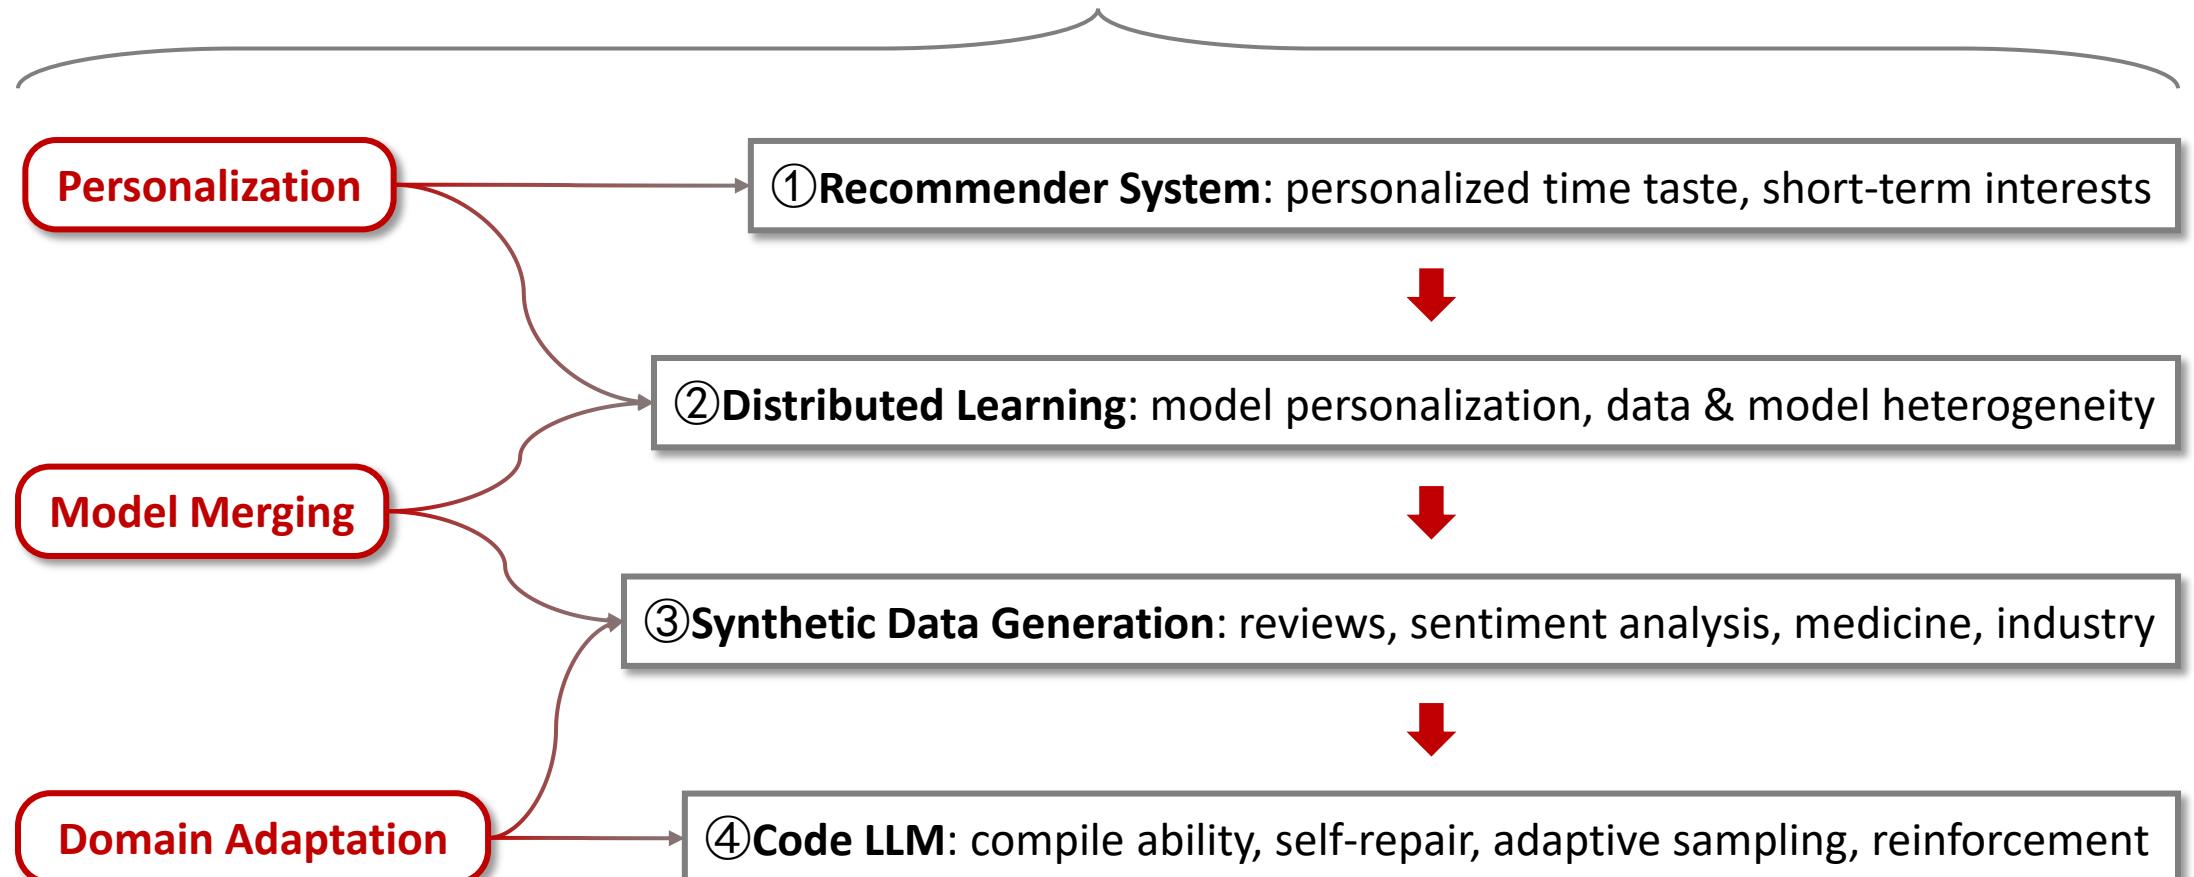
# 总览

- **研究兴趣:** 垂域自适应(核心)、模型融合、模型个性化
- **研究领域(多领域迁移能力):** 代码大模型、合成数据集生成、分布式机器学习、推荐系统
- **突出优势:** 科研能力 + 工程经验
- **论文发表:** 9 篇一作顶会顶刊论文
  - 阶段① AAAI'23 (oral), KDD'23, ICCV'23, NeurIPS'23, JMLR'25
  - 阶段② AAAI'24, CVPR'24, KDD'25
  - 阶段③ EMNLP'24, ICML'25, ICML'25 (spotlight)
- **开源项目(发起人):**
  - EvolveGen, PFLlib (**1800+** stars, **300+** forks), HtFLlib, HtFLlib on Device, FL-IoT, etc.
- **获奖:** 中国科协-博士青年人才托举(中国人工智能协会), 吴文俊人工智能荣誉博士, 博士生国家奖学金
- **落地项目:** ①与医院合作进行跨医院癌症相关研究、②为12345政务服务热线智能模型进行跨省份知识迁移、③在20+单片机上部署异构模型分布式训练、④带9人团队搭建分布式机器学习平台, 交付给数据中心
- **影响力:** **700+** 谷歌学术引用, **30K+** 主流媒体曝光, 受到IEEE/ACM Fellows在CCF大会上对我工作的赞扬
- **实习交流:** 字节跳动 AML, 清华智能产业研究院, 阿卜杜拉国王科技大学 SANDs lab, 腾讯代码模型组

# Systematical Research Trace



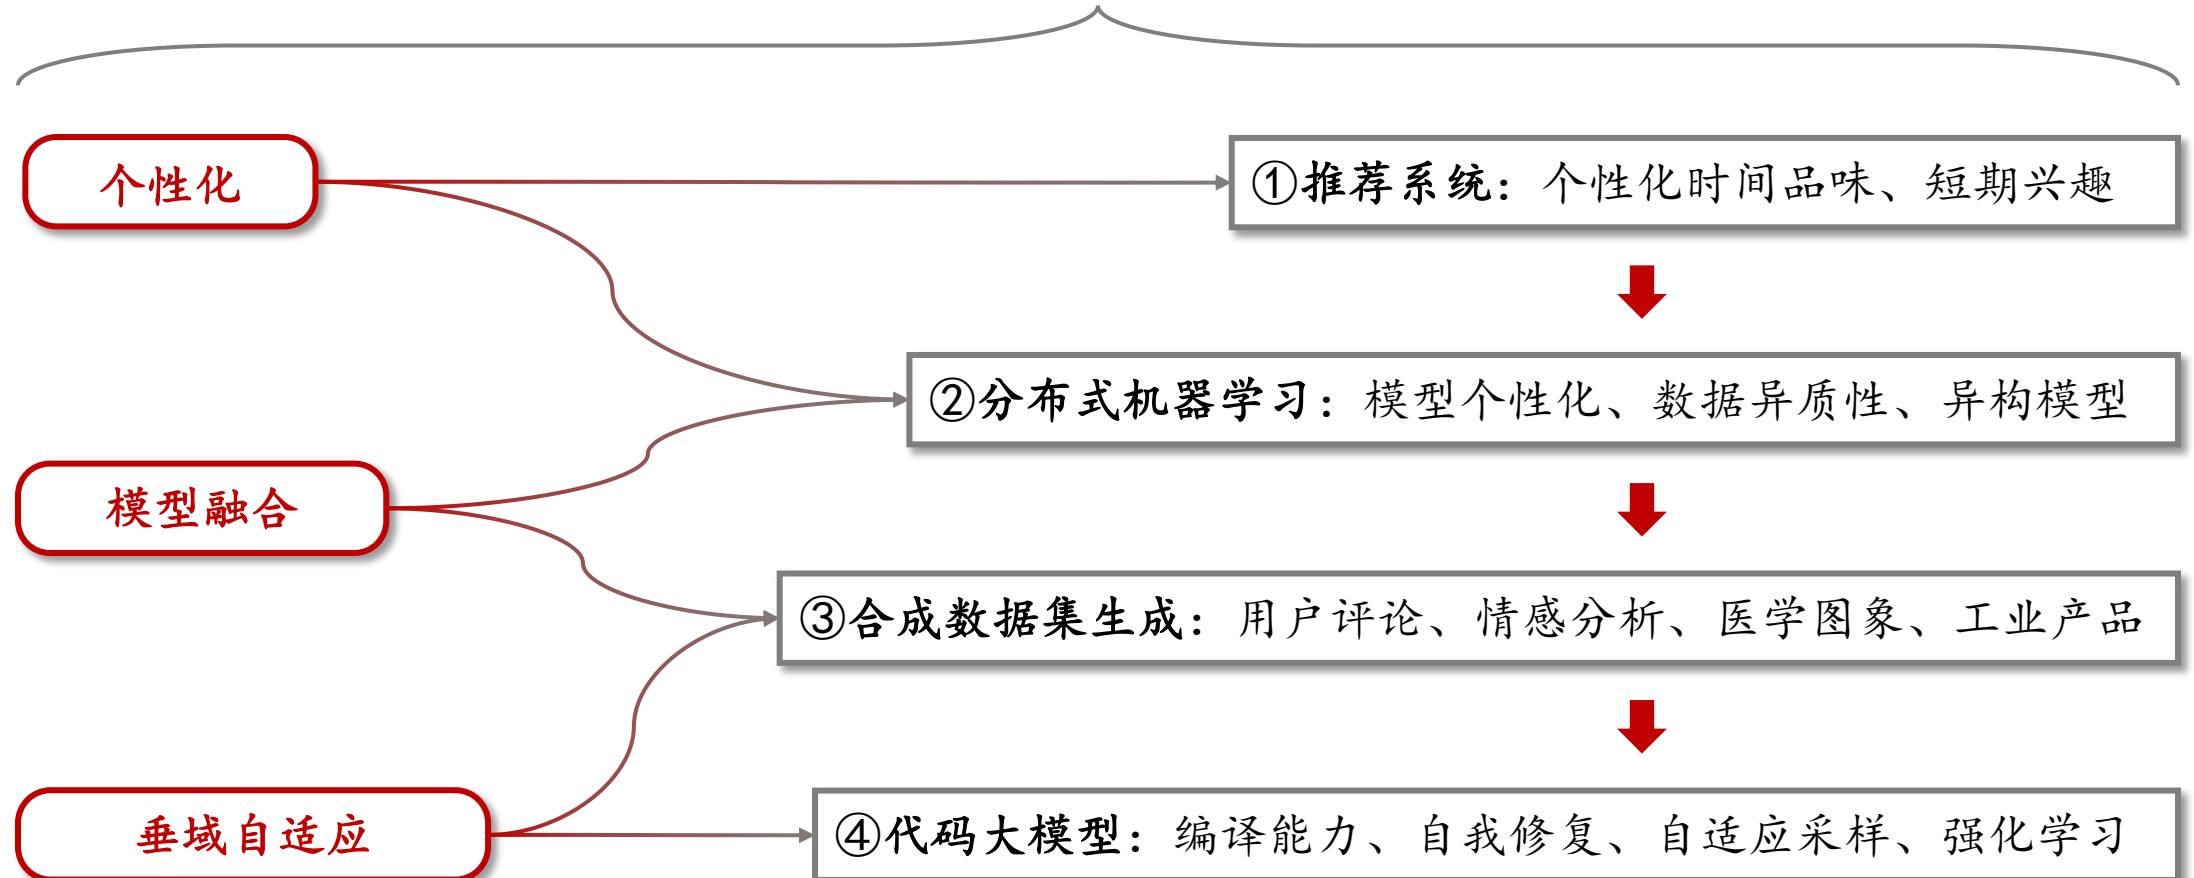
**Scientific Problem:** Balancing generalization and specialization in special domains





# 系统性科研路径

科学问题：平衡特定领域中模型的泛化能力和专业化能力





# Code LLM (Code)

- **CodeBuddy**, Cursor, Claude Code, GitHub Copilot, Trae, Lingma, CodeFuse, etc.

The screenshot displays the Code LLM (Code) platform interface, featuring several key components:

- Header:** Includes a navigation bar with links to 首页 (Home), 安装教程 (Installation Guide), 产品价格 (Product Price), 在线对话 (Online Chat), 企业开通 (Enterprise Activation), 帮助文档 (Help Document), API 文档 (API Document), and AI IDE.
- Left Sidebar:** Features a logo, a "登录" (Login) button, and sections for "你好, 我是 CodeBuddy" (Hello, I am CodeBuddy), "技术问题查询" (Technical Problem Inquiry), "知识库深度检索" (Knowledge Base Deep Search), and "安装到你的IDE中, 大幅提升你的编码效率!" (Install into your IDE, significantly improve your coding efficiency!). It also includes a search bar with placeholder text "请输入你想查询的内容".
- Middle Content Area:** Contains several cards:
  - MCP: 支持外部工具调用**: Describes MCP's support for external tool invocation, mentioning GitHub, Firecrawl, Apidog, and Playwright.
  - 代码补全 Plus**: Shows a snippet of code for creating a SCF client using the TencentCloud SDK.
  - 工程理解智能体 Plus**: Discusses AI's role in understanding project engineering.
  - 智能问答**: Shows a question about the principle of two-way data binding and an answer from CodeBuddy.
  - 代码评审**: Shows a code review interface with tabs for Craft, Chat, Code Review (selected), and Unit test, displaying findings for Detail.vue.
- Bottom Buttons:** Includes links for JetBrains and VS Code integration.



# Synthetic Data Generation (SDG)

- Given a **prompt**, with or without **data examples**,
- AI generates a dataset that **aligns with the user's request**.
  - Focusing on special domains (e.g., code, medicine, industry, etc.)

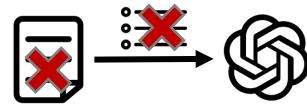


# [SDG]: PCEvolve (Domain Adaptation)

- Widely-used approaches in specialized domains for large models:
  - Fine-tuning
    - **Costly** for large model training, **data scarcity**
  - Few-shot in-context learning (ICL)
    - **Privacy issue**, effortful **prompt engineering**
  - Zero-shot ICL + selection
    - **Costly** for large amount data generating, effortful **prompt engineering**



Fine-tuning



Few-shot ICL

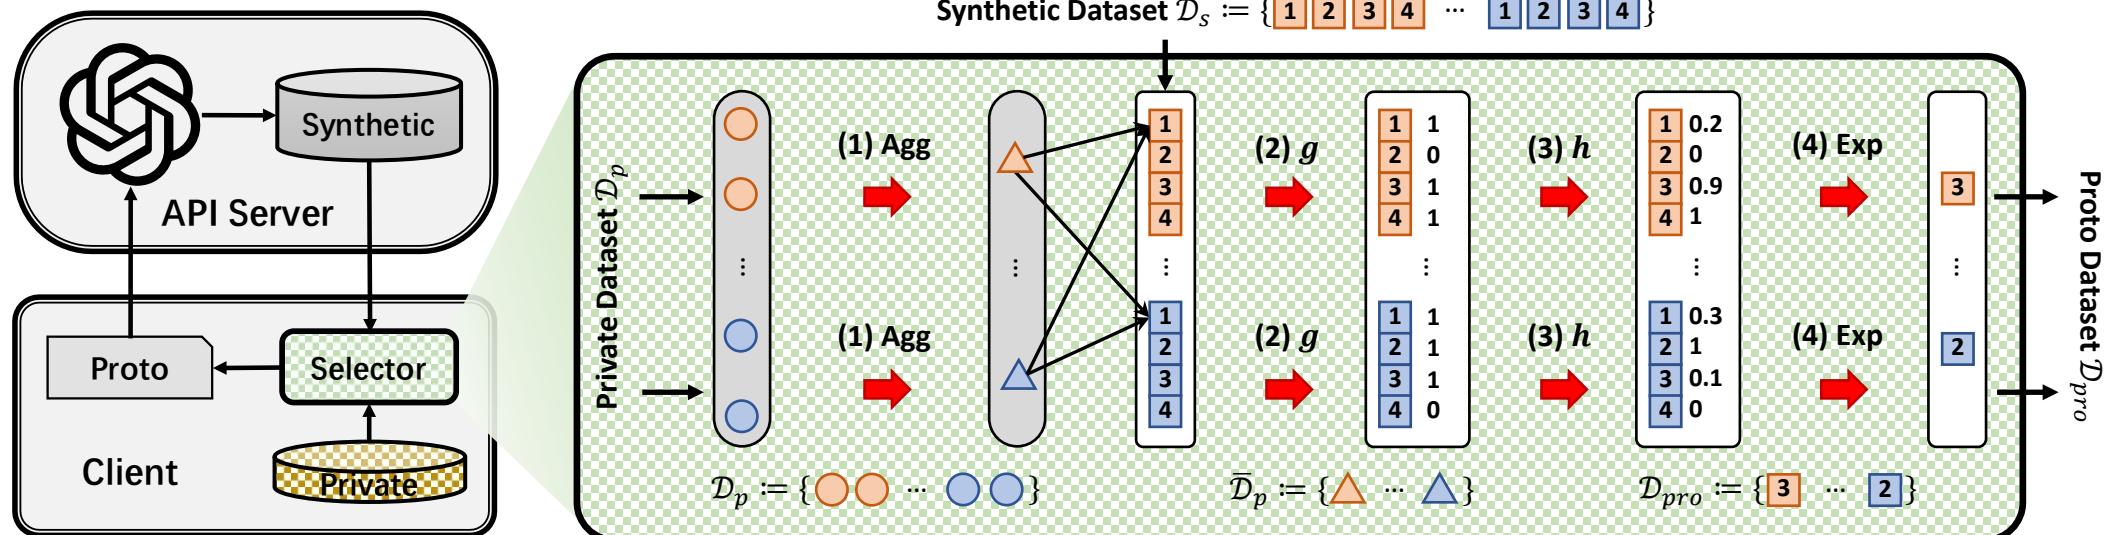


Zero-shot ICL



# [SDG]: PCEvolve (Domain Adaptation)

- Via iterative evolution of synthetic datasets, **you only need to provide a few labeled samples** — we'll **evolve** an entire dataset for you,
- While **protecting privacy**





# [SDG]: PCEvolve (Domain Adaptation)

- **COVIDx**: chest X-ray images for COVID-19
- **Came17**: tumor tissue patches from breast cancer metastases
- **KVASIR-f**: endoscopic images for gastrointestinal abnormal findings detection
- **MVAD-I**: leather surface anomaly detection

Top-1 accuracy (%) on four specialized datasets

	COVIDx	Came17	KVASIR-f	MVAD-I
Init	49.34	50.47	33.43	33.33
RF	50.01	54.82	34.66	48.17
GCap	50.86	55.77	32.66	27.33
B	50.42	54.41	32.57	43.21
LE	50.02	55.44	35.51	27.93
DPImg	49.14	61.06	33.35	37.03
PE	59.63	63.66	48.88	57.41
PE-EM	57.60	63.34	43.01	50.06
PCEvolve-GM	56.91	62.63	43.55	55.56
PCEvolve	<b>64.04</b>	<b>69.10</b>	<b>50.95</b>	<b>59.26</b>

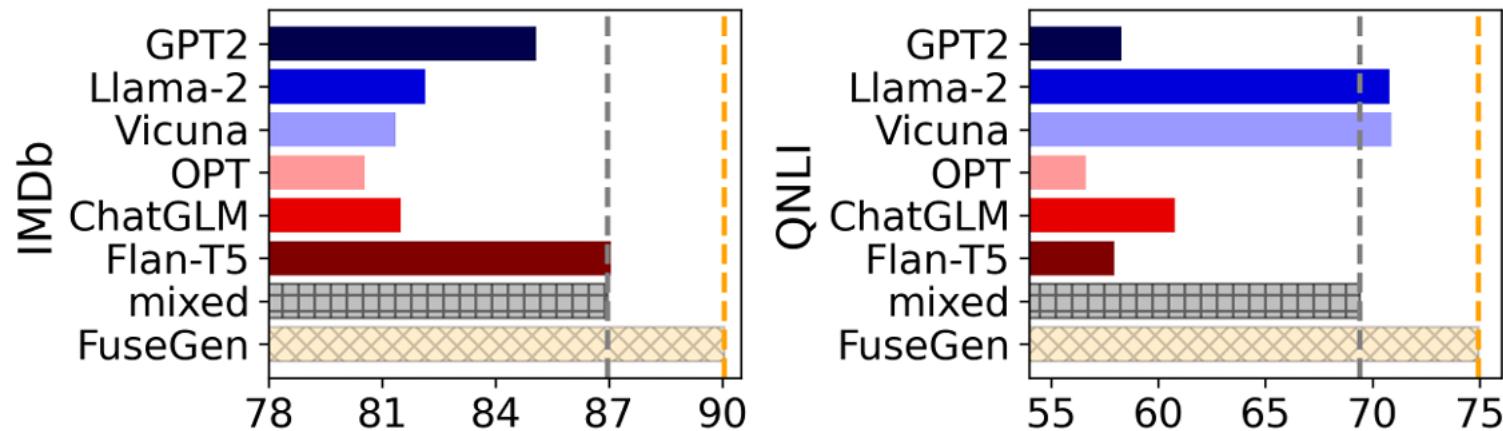
# [SDG]: PCEvolve (Domain Adaptation)



Generated leather surface images w.r.t. MVAD-I for industry anomaly detection. The three rows show normal images, cut defects, and droplet defects. “Initial” denotes API-generated images using just the prompt. “Private” denotes the real images from MVAD-I.

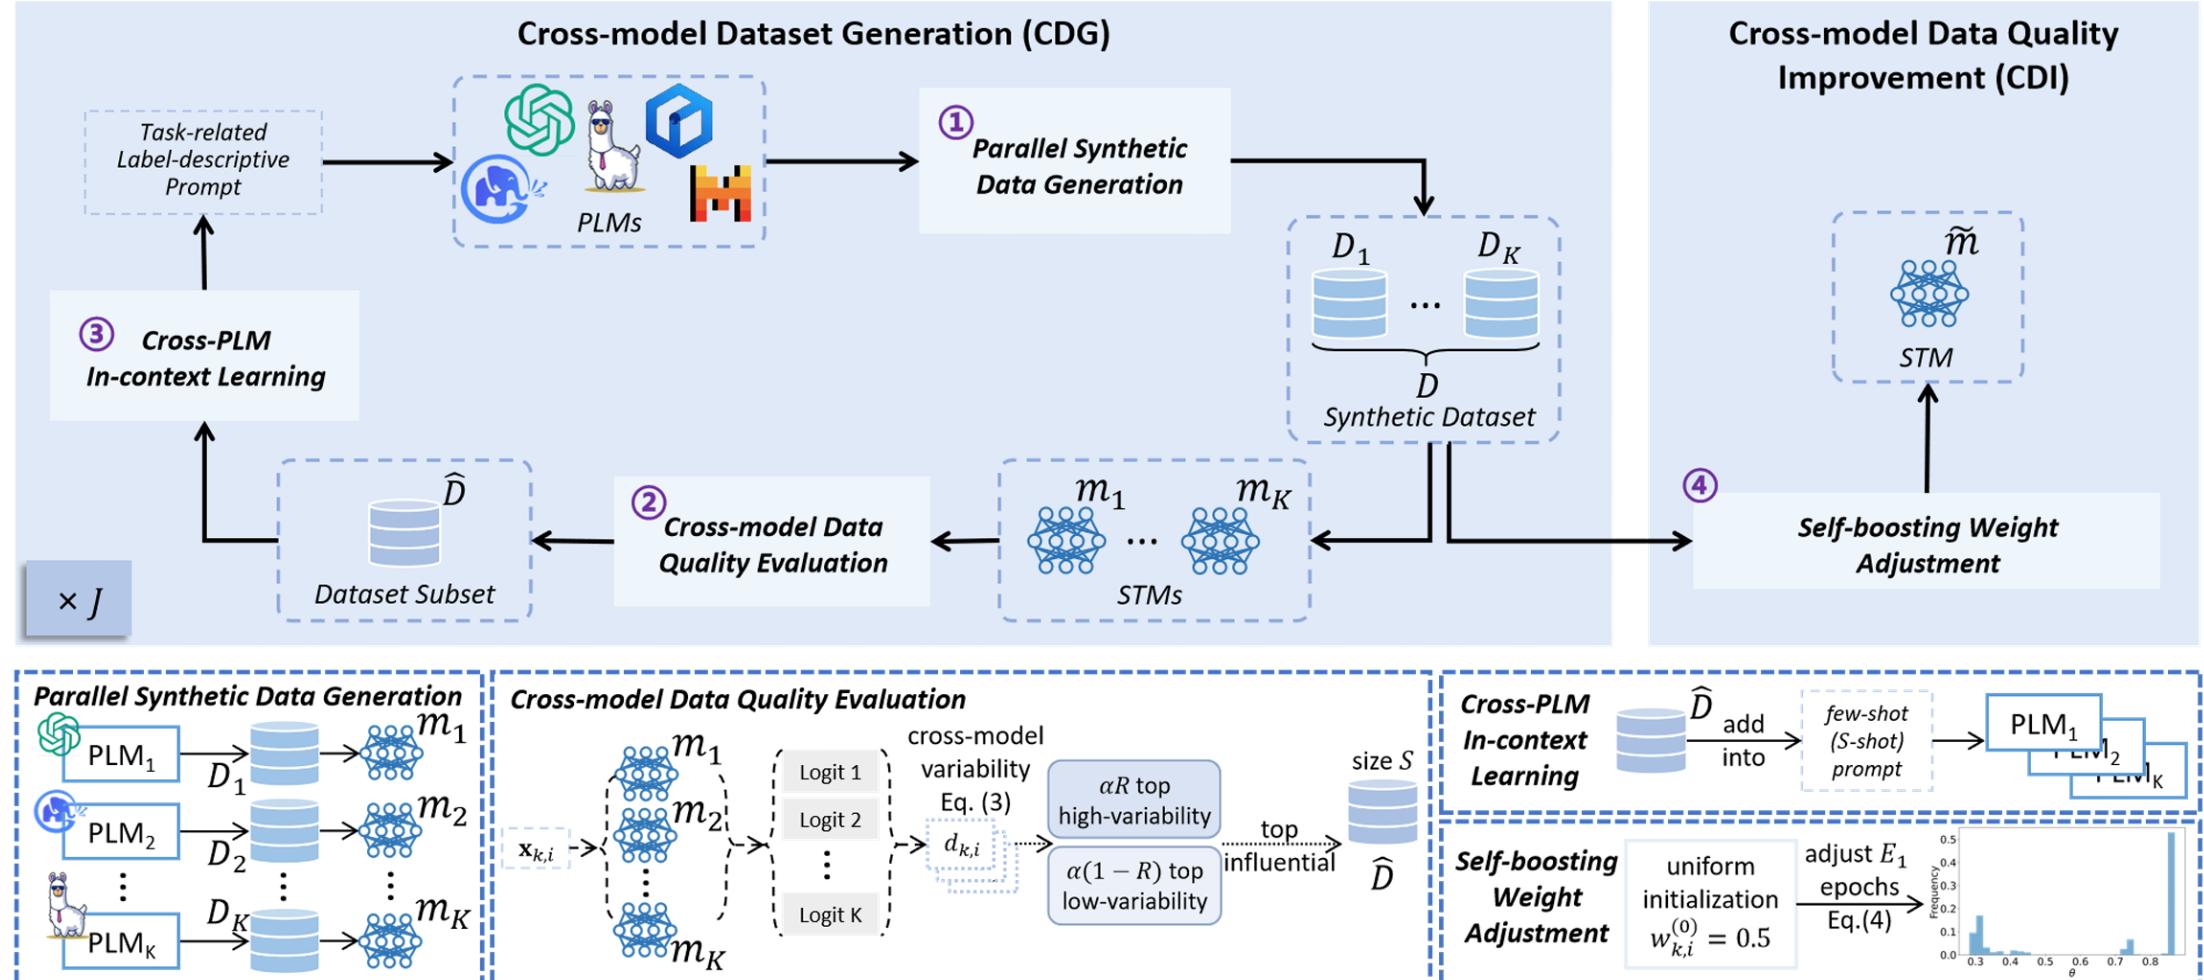
# [SDG]: FuseGen (Domain Adaptation, Model Merging)

- Pre-trained Language Models (PLMs) have **different tastes**
- Merging their outputs to create **diverse datasets**,
- Through **evolution**



# [SDG]: FuseGen (Domain Adaptation, Model Merging)

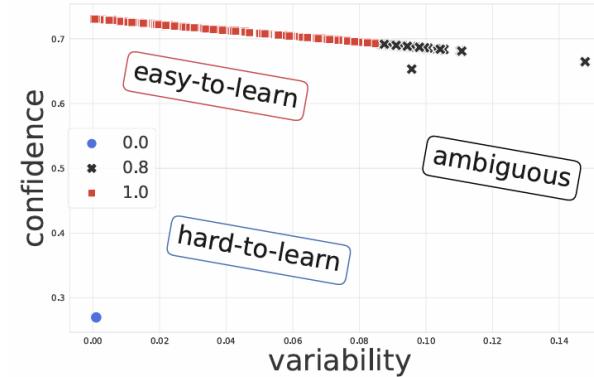
- We consider downstream models' feedback as reward signals for evolution



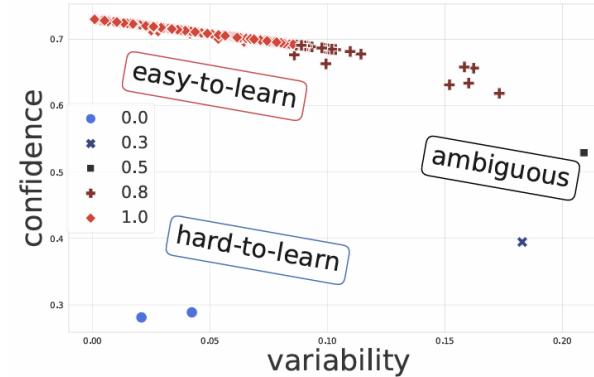
# [SDG]: FuseGen (Domain Adaptation, Model Merging)



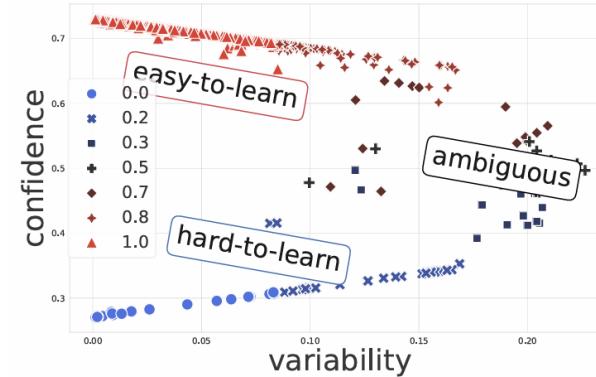
- Synthetic dataset cartography



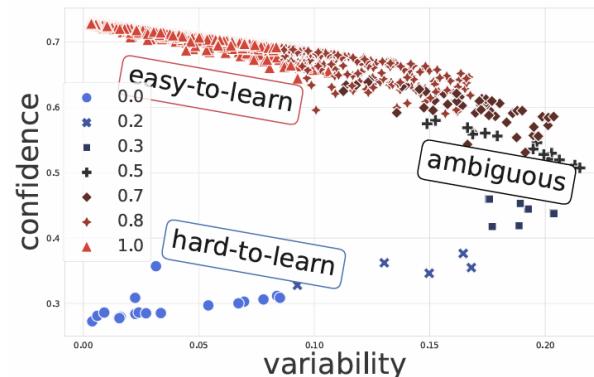
(a) Llama-2 ZeroGen  $K = 1$  (84.23)



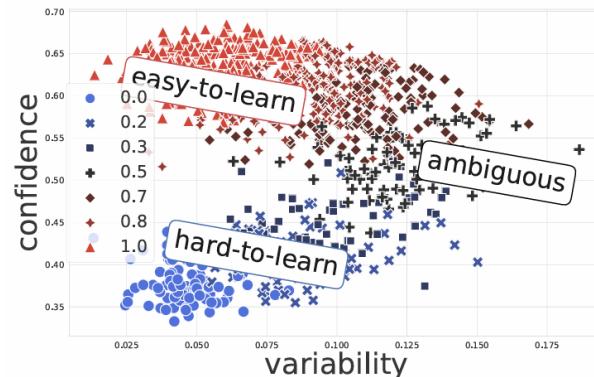
(b) Llama-2 ProGen  $K = 1$  (84.24)



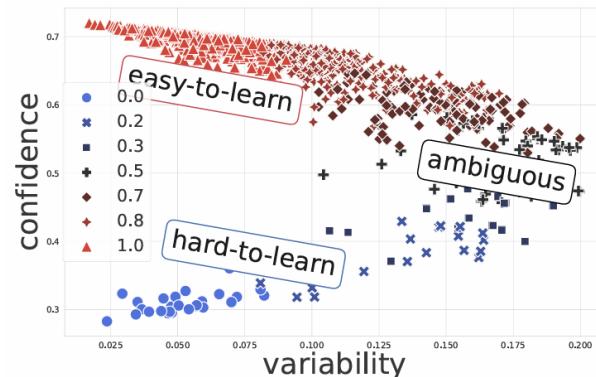
(c) Llama-2 Ours  $K = 6$  (86.60)



(d) Flan-T5 ZeroGen  $K = 1$  (88.18)



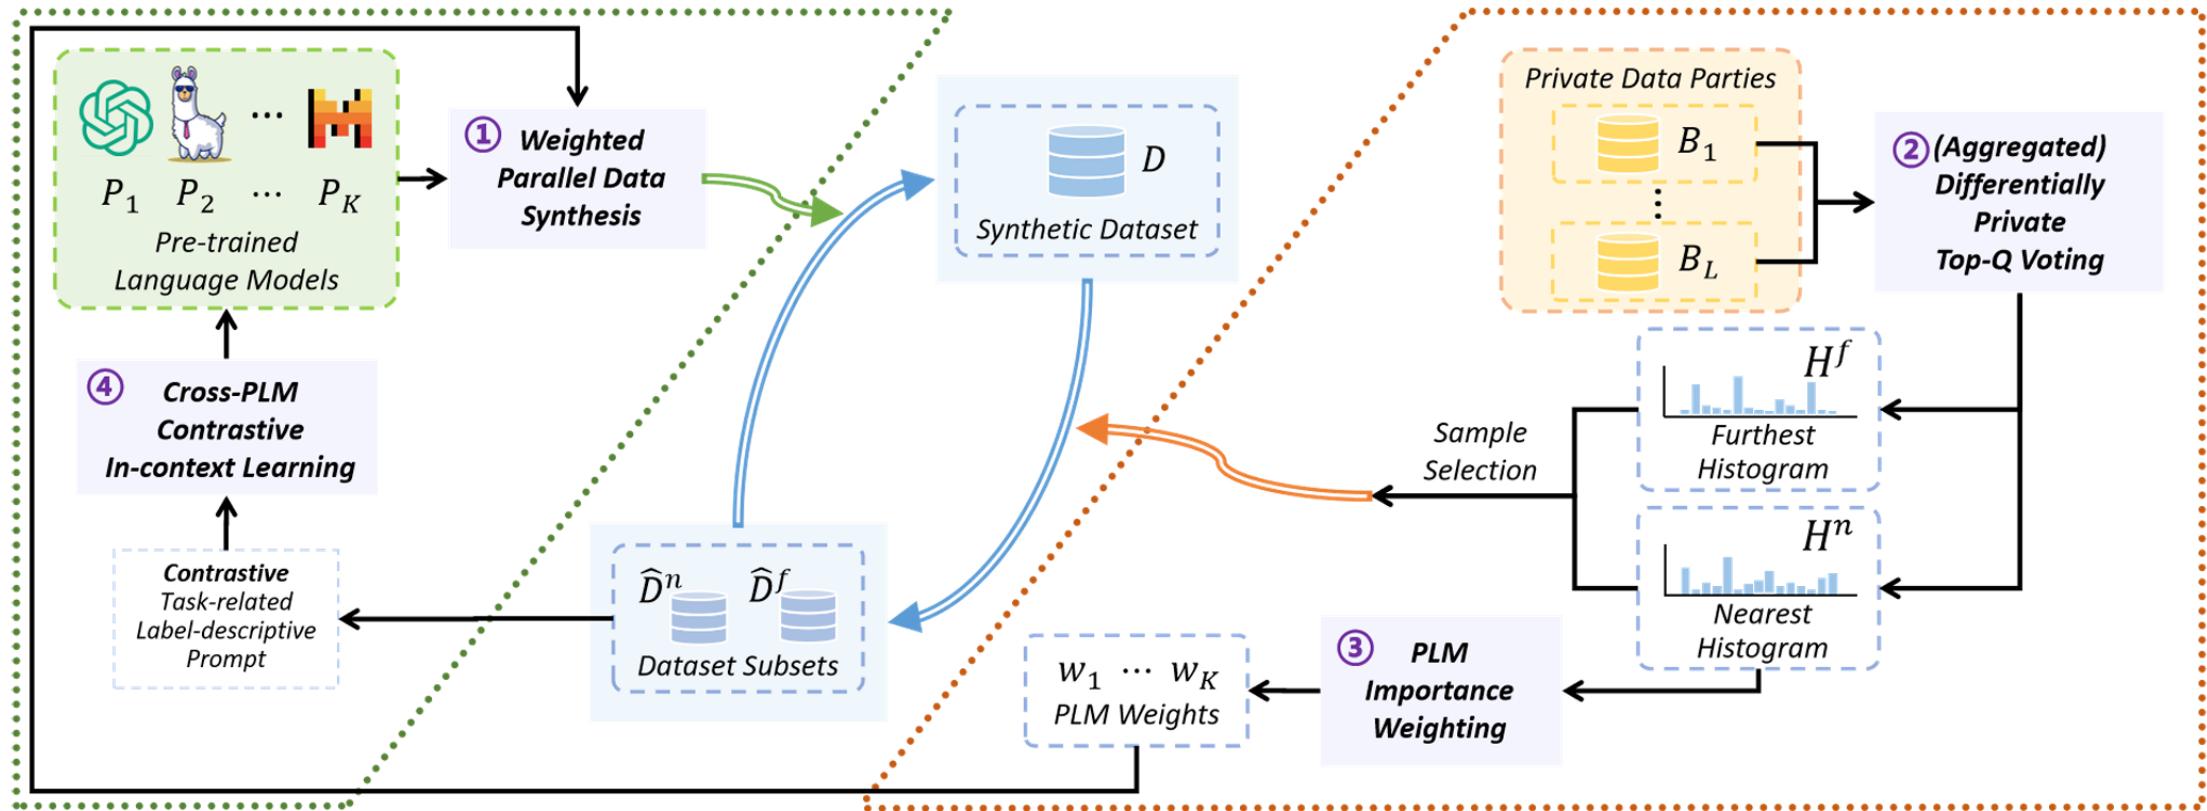
(e) Flan-T5 ProGen  $K = 1$  (85.80)



(f) Flan-T5 Ours  $K = 6$  (88.73)

# [SDG]: WASP (Domain Adaptation, Model Merging)

- Users provide **diverse private data samples**, PLMs generate **diverse datasets**
- **Contrastive voting-based sample selection for evolution**



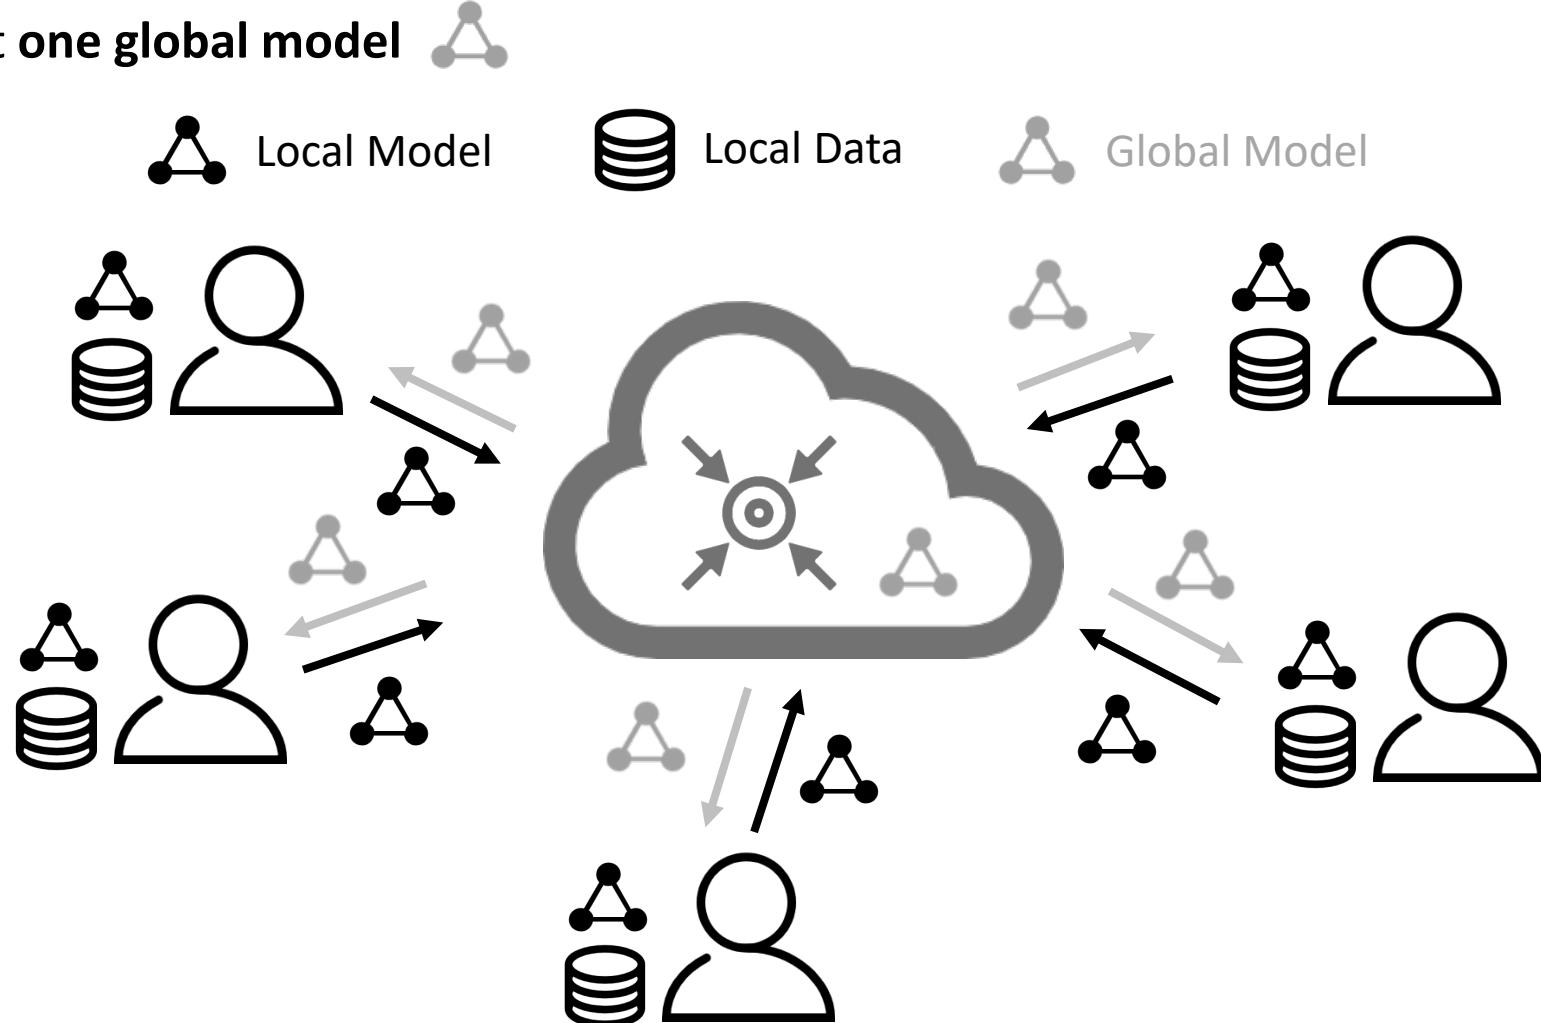


# [SDG]: WASP (Domain Adaptation, Model Merging)

Model	High-quality Sample Text (Examples)	Label
GPT-2	"This movie isn't for everyone, but it's worth a look. But if you have to be in the dark to enjoy it, just buy the soundtrack."	positive
GPT-2	"is that it has some semi-interesting world-building, but the information it offers is grossly insufficient to allow one to progress in the story. The central characters have little to do with the world beyond playing around with their toys and talking to themselves."	negative
Llama-2	"Incredible film! I laughed, I cried, and I left the theater feeling inspired. Every performance was top-notch and the direction was flawless. A true masterpiece that will be remembered for years to come. 10/10."	positive
Llama-2	"This film was a colossal disappointment. I found myself glancing at my wrist every 10 minutes, eagerly anticipating its conclusion. The narrative was intricate and challenging to comprehend, and the personas were underdeveloped. I had no emotional investment in their plight. The performances were lacking, and the dialogue was trite and unconvincing. In summary, I would not advise viewing this motion picture to any individual."	negative
Vicuna	"A movie for people who want to see an action-packed adventure that will keep them on the edge of their seats."	positive
Vicuna	"Felt like a missed opportunity. Glacial pacing and shallow characters. The premise was intriguing but the execution fell flat. Avoid this one."	negative
OPT	"-This film is enjoyable to watch because of its well-developed storyline and most of all its storyline that i can relate to and would like to have. This film also has good acting performances and the songs help the film a lot. It's not a typical love story kind of a film, but it does have lots of interesting messages to give. Very well-done! Definitely recommend this film!"	positive
OPT	"Quote: The first sentence means nothing to me. Without context the first two sentences also mean nothing."	negative
ChatGLM3	"Attention getter, visually interesting and outstanding acting, the story of an American citizen that is in Mexico and gets involved in a murder is a good movie."	positive
ChatGLM3	"This model is an instruction-following model off-the-shelf trained on the Stanford Large Scale US Air Force data set."	positive

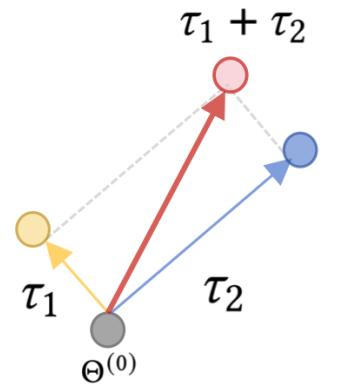
# Federated Learning (FL) (Model Merging, Personalization)

- A **collaborative** and **privacy-preserving** technique for AI model training
- Finally output **one global model**

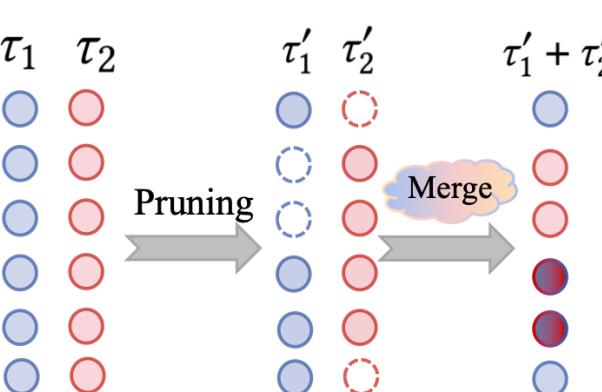
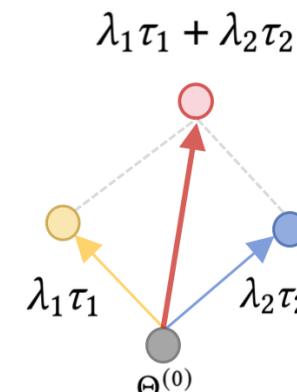


# Federated Learning (FL) (Model Merging, Personalization)

- I mainly focus on **model merging** in FL, which is also **popular** in **large model training** by
  - merging parameters, merging intermediate features, merge parameter-efficient LoRA modules, etc.
- Multi-modal model:** obtain a single, effective, and parameter-efficient modality-agnostic model
- RL:** DogeRM [1] merges the reward model with LLMs fine-tuned on different downstream domains to create domain-private reward models directly



(a) Weighted-based Merging



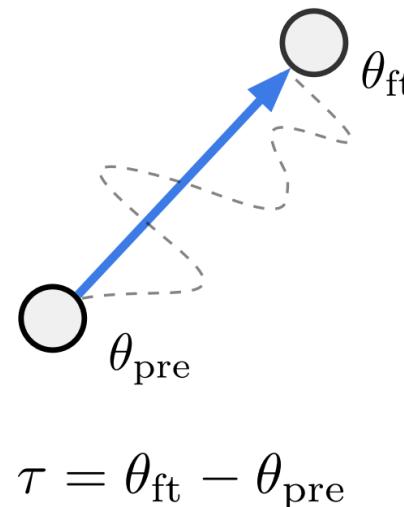
(b) Subspace-based Merging

(c) Routing-based Merging

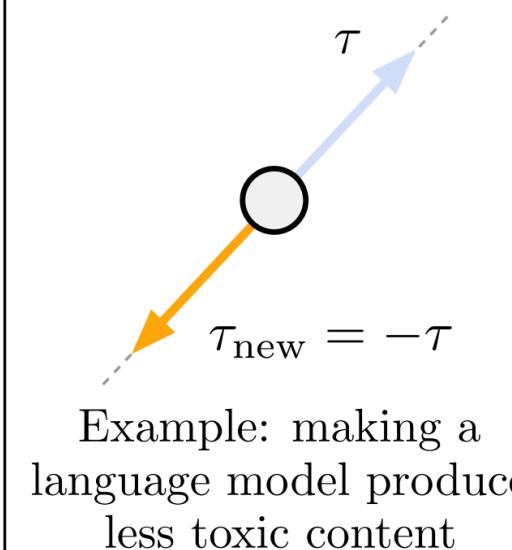
# Federated Learning (FL) (Model Merging, Personalization)

- I mainly focus on **model merging** in FL, which is also **popular** in **large model training** by
  - merging parameters, merging intermediate features, merge parameter-efficient LoRA modules, etc.
- Model editing:** [1] shows that task vectors can be modified and combined together, and the behavior of the resulting model is steered accordingly. AlphaEdit [2] (ICLR'25 best paper)

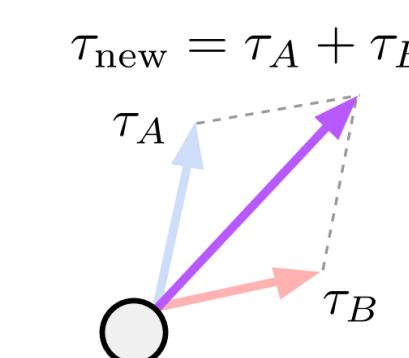
a) Task vectors



b) Forgetting via negation

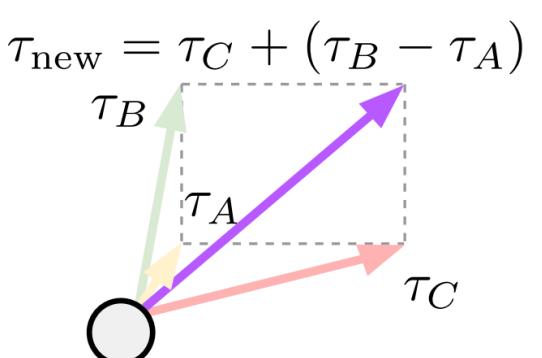


c) Learning via addition



Example: building a multi-task model

d) Task analogies



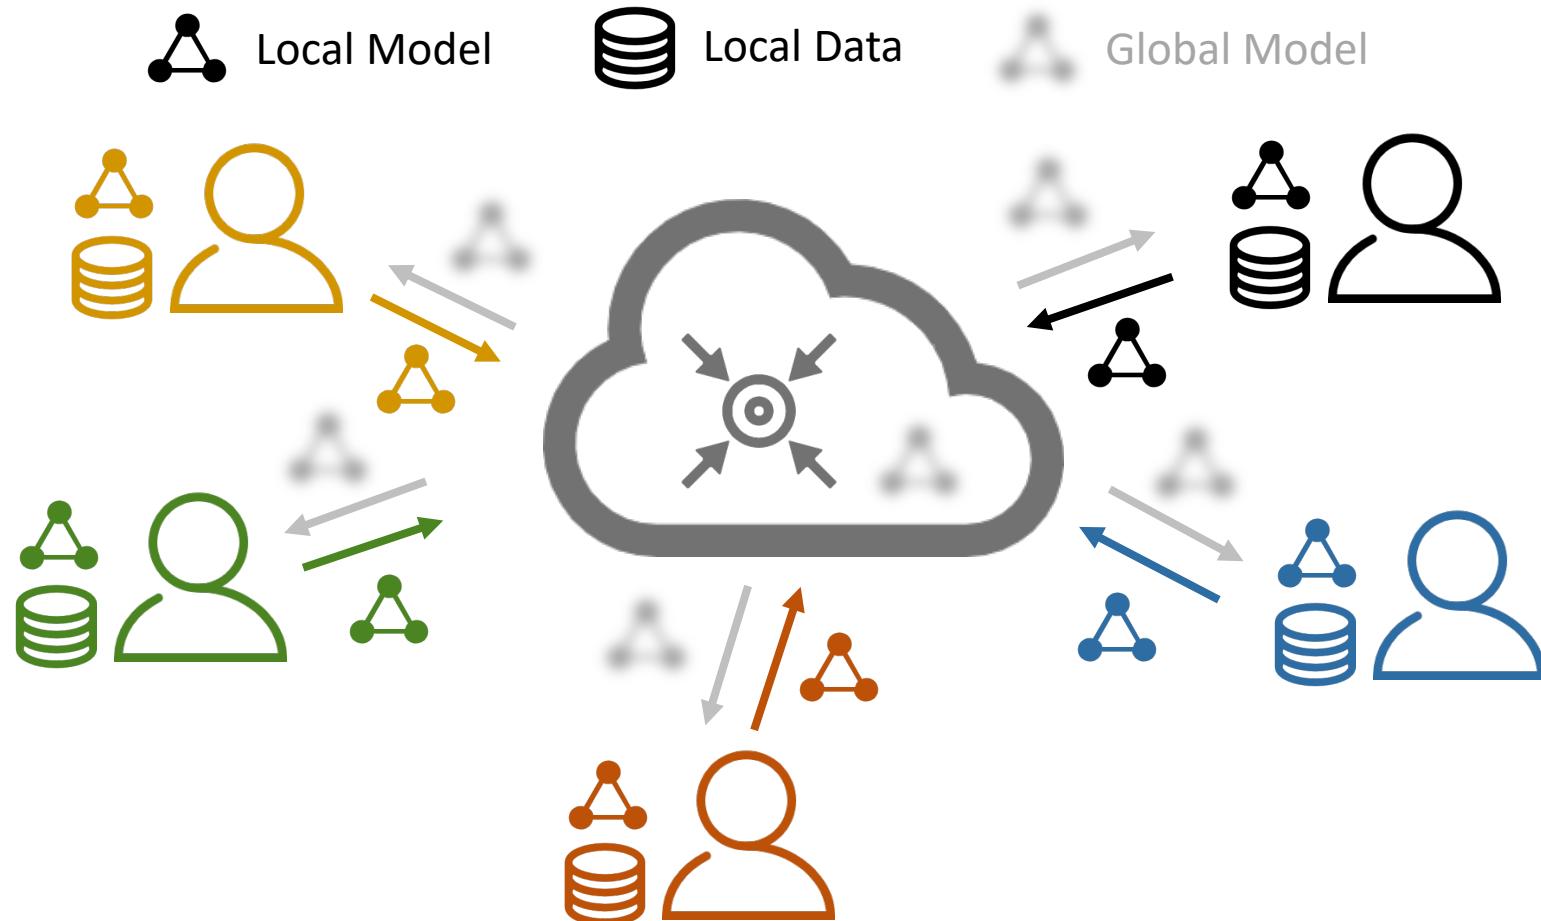
Example: improving domain generalization

[1] Ilharco, Gabriel, et al. "Editing models with task arithmetic." *ICLR* 2023.

[2] Fang, Junfeng, et al. "Alphaedit: Null-space constrained knowledge editing for language models." *ICLR* 2025.

# [FL]: Data Heterogeneity (Model Merging, Personalization)

- Different tasks have **different** data distributions
- **Personalized federated learning (pFL)** comes along



# [FL]: PFLlib, pFL algorithm library and benchmark

- Beginner-friendly
- 39 FL&pFL, 3 scenarios, 24 datasets
- Popular (1800+ stars)
- 500 clients: 5GB GPU memory
- Rapidly developing:

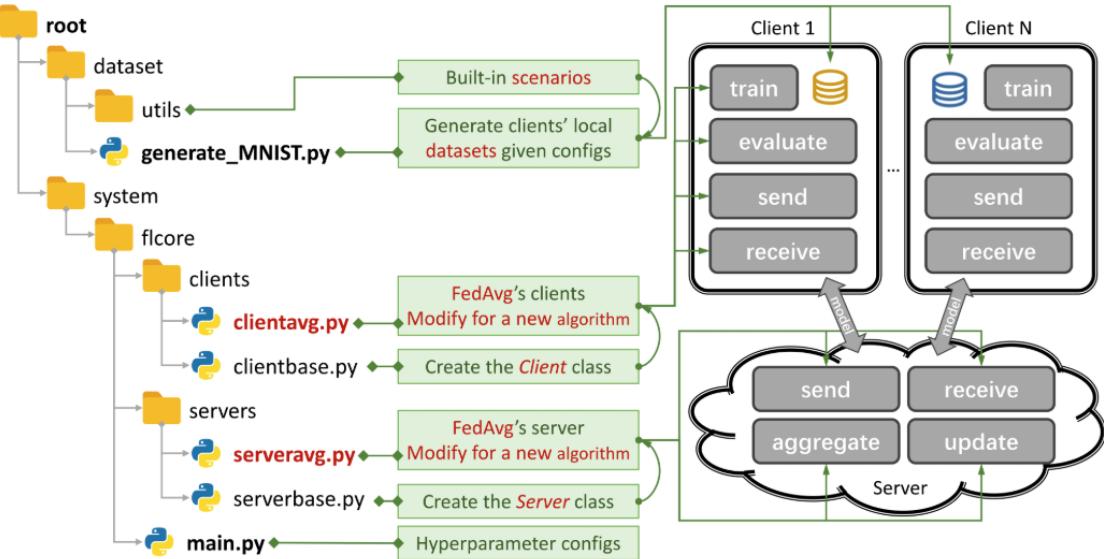
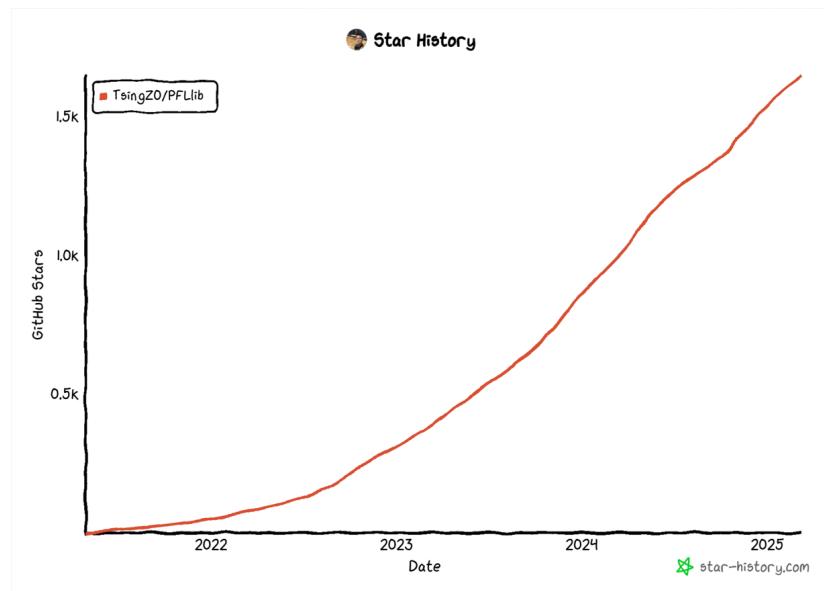


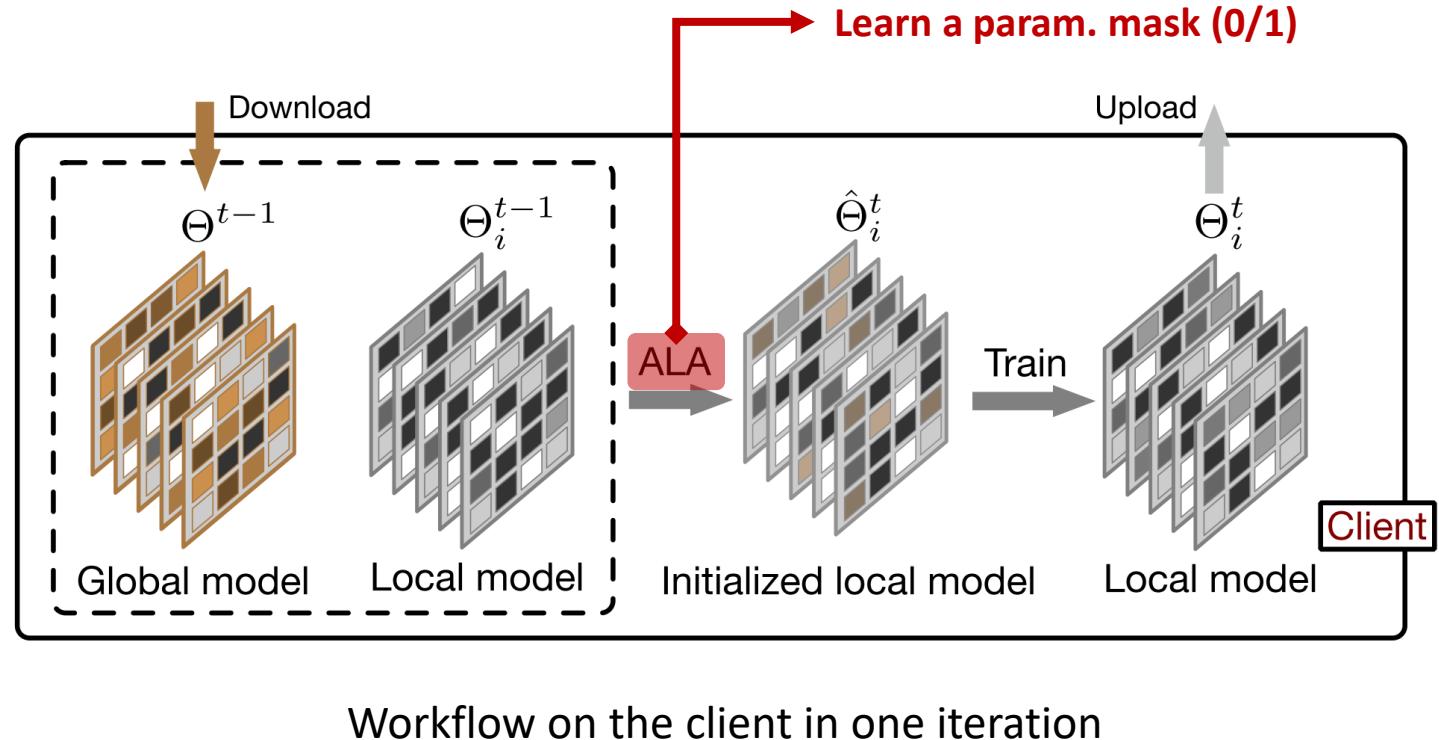
Figure 1: An Example for FedAvg. You can create a scenario using `generate_DATA.py` and run an algorithm using `main.py`, `clientNAME.py`, and `serverNAME.py`. For a new algorithm, you only need to add new features in `clientNAME.py` and `serverNAME.py`.

The screenshot shows the PFLlib website with two main sections:

- Benchmark Platform:** Describes the platform for comparing algorithms across various datasets and scenarios.
- Leaderboard:** Displays the test accuracy (%) on CV and NLP tasks for different methods (FedAvg, FedProx, FedDNN, etc.) under different label skew settings (Pathological Label Skew Setting and Practical Label Skew Setting).

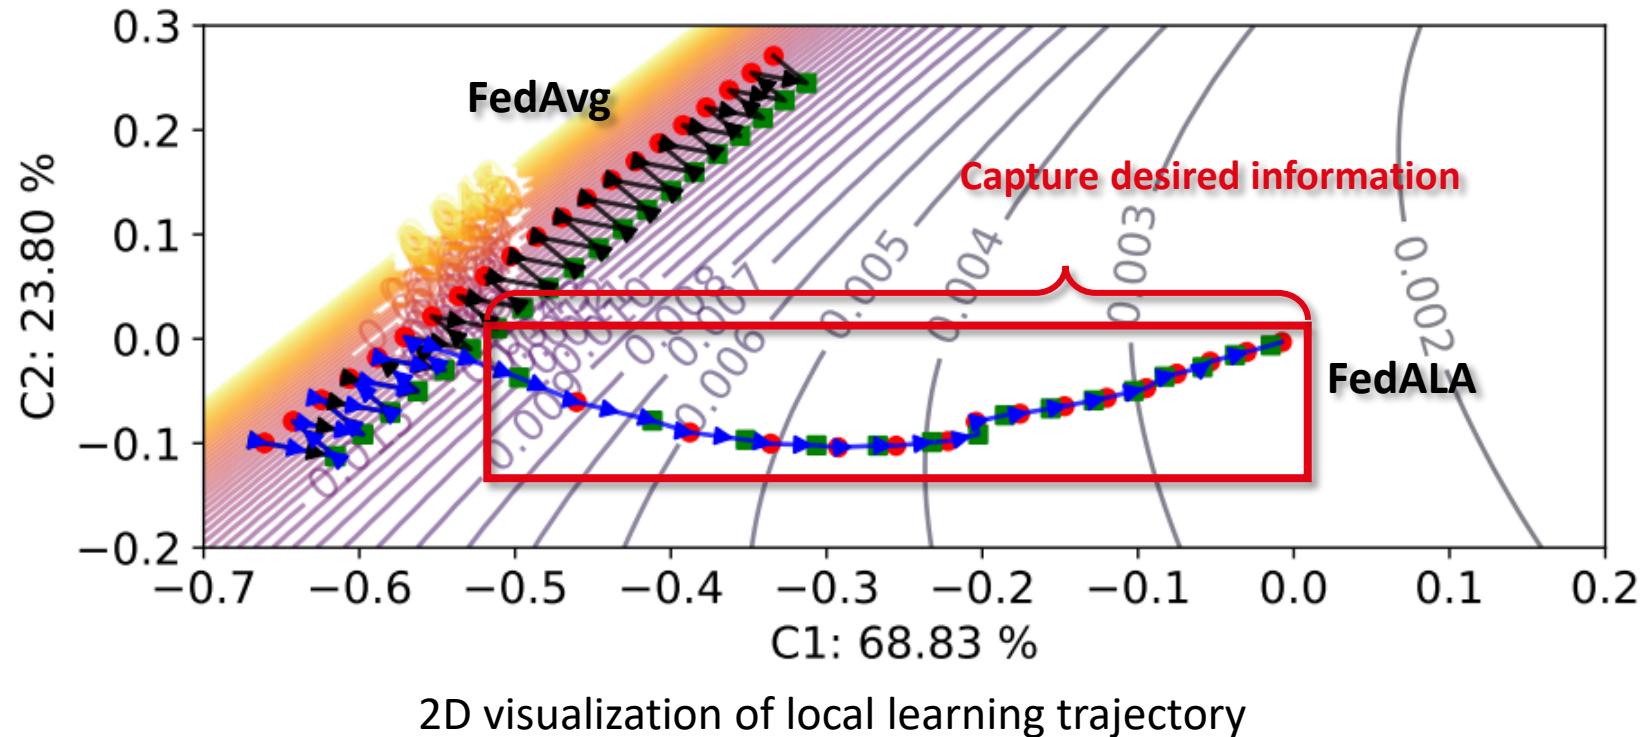
# [FL]: FedALA (Model Merging, Personalization)

- Extract each client's desired information from the **global model** that facilitates local training
- Adaptively aggregate the information in the global and local model for initialization



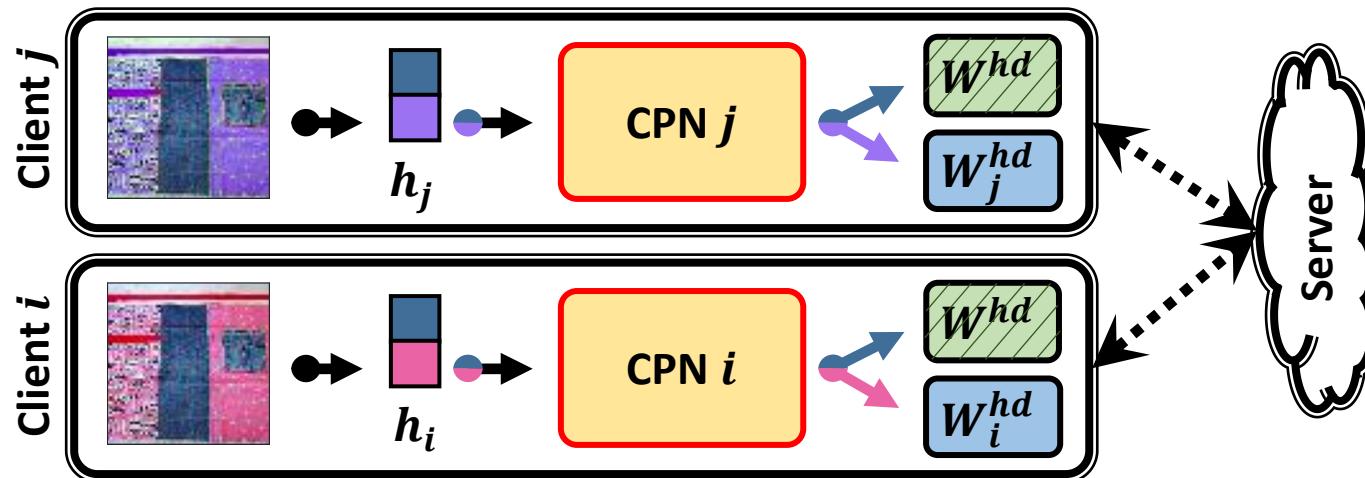
# [FL]: FedALA (Model Merging, Personalization)

- Learning trajectory on one client: **FedAvg** vs. **FedALA**
- Activate ALA in the subsequent iterations



# [FL]: FedCP (Model Routing, Personalization)

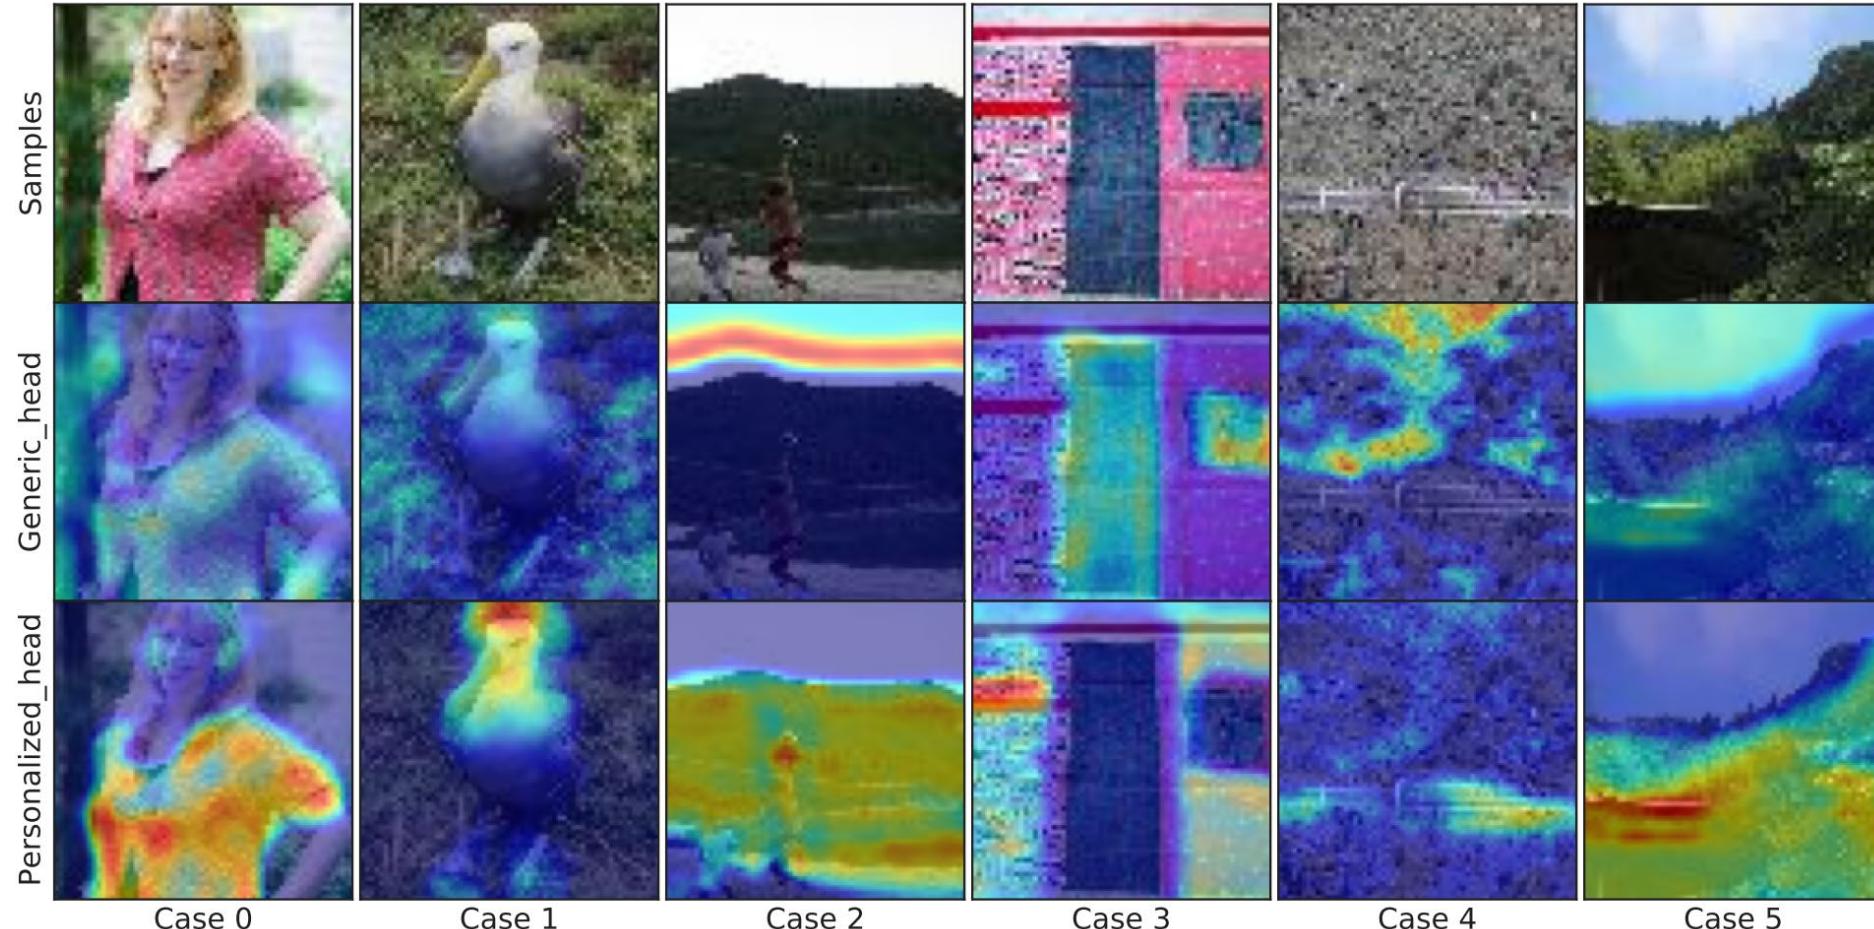
- We separate feature information via an auxiliary **Conditional Policy Network (CPN)**.
  - Sample-specific separation
  - Lightweight (e.g., 4.67% parameters of ResNet-18)



- Then, we utilize global and personalized information via global and personalized heads.

# [FL]: FedCP (Model Routing, Personalization)

- Separating Feature Information



Six samples from the Tiny-ImageNet dataset

# [FL]: GPFL (Model Routing, Personalization)

- GCE introduces more global information **simultaneously** with local training
- CoV **eliminates the interaction** between global and personalized feature learning

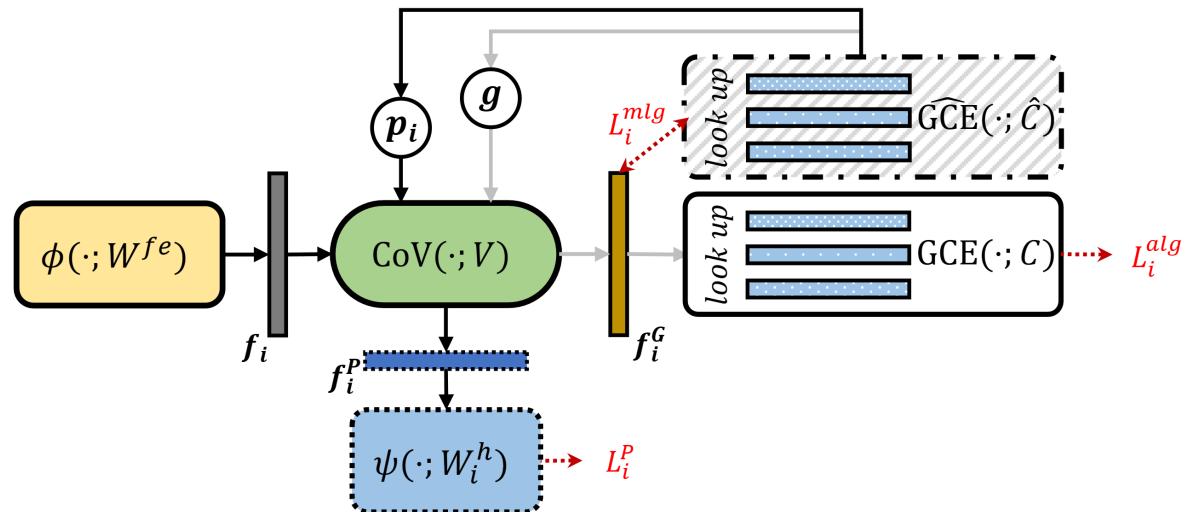
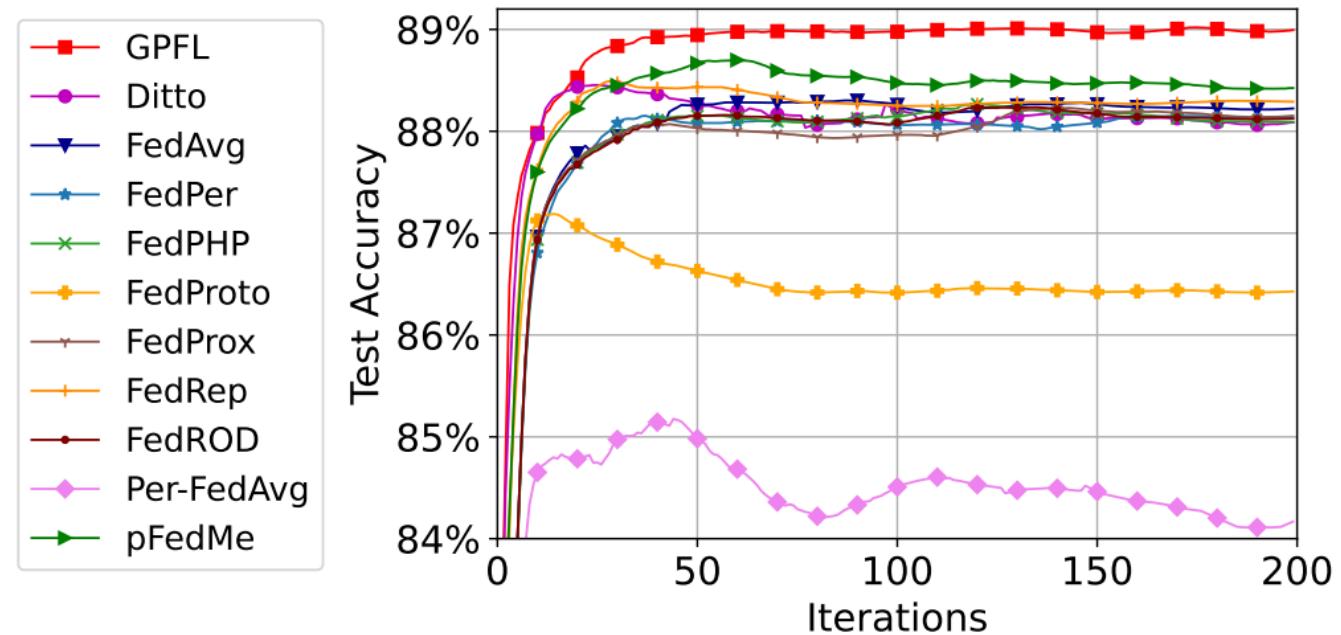


Illustration of client modules and data flow between them

# [FL]: GPFL (Model Routing, Personalization)

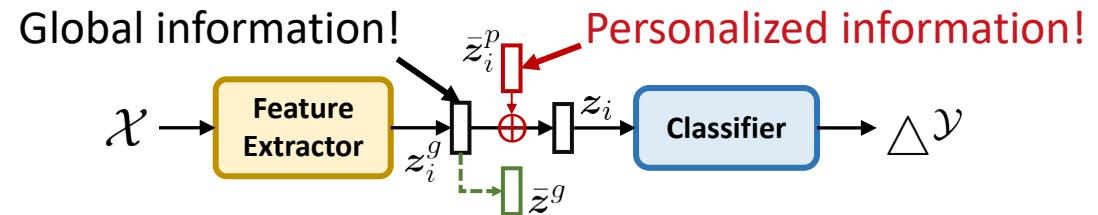
- Relieve the **widely existed** overfitting issue in pFL



Test accuracy curves in the feature shift setting

# [FL]: DBE (Feature Decoupling, Personalization)

- Eliminate domain bias by store **personalized information** in PRBM
- Enhance **information disentanglement** by guiding feature extractor with MR



Local model (with PRBM and MR)



# [FL]: DBE (Feature Decoupling, Personalization)

- Improve bi-directional knowledge transfer

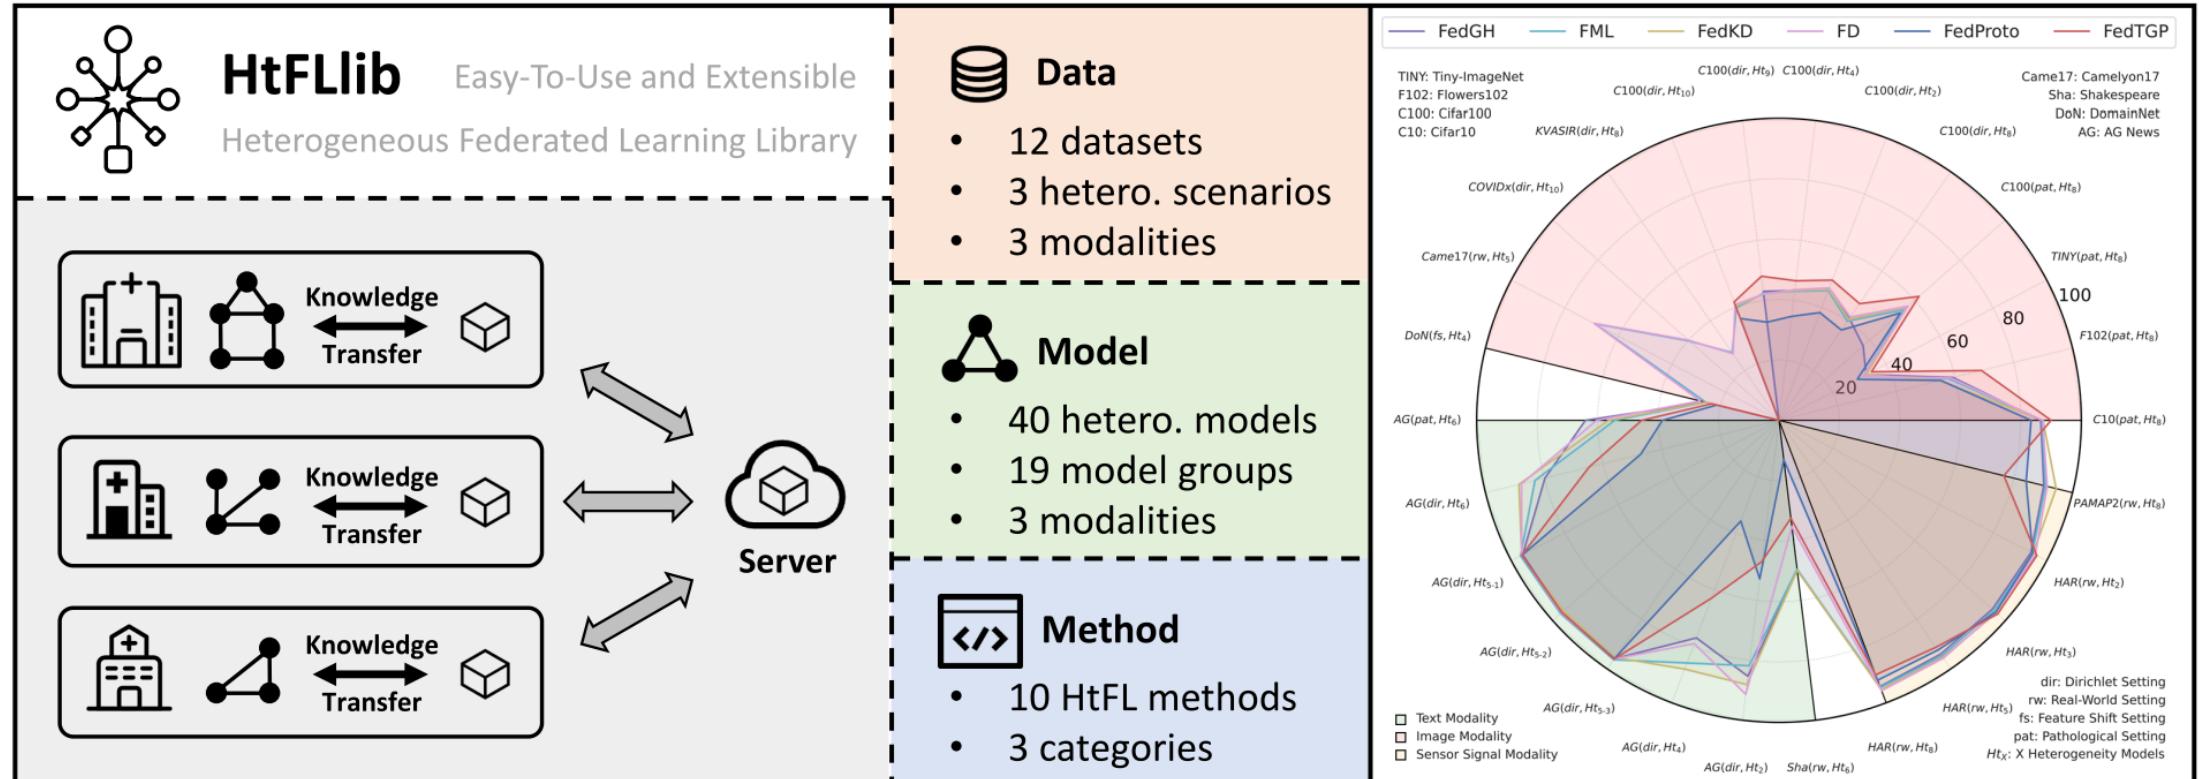
**Corollary 1.** Consider a local data domain  $\mathcal{D}_i$  and a virtual global data domain  $\mathcal{D}$  for client  $i$  and the server, respectively. Let  $\mathcal{D}_i = \langle \mathcal{U}_i, c^* \rangle$  and  $\mathcal{D} = \langle \mathcal{U}, c^* \rangle$ , where  $c^* : \mathcal{X} \mapsto \mathcal{Y}$  is a ground-truth labeling function. Let  $\mathcal{H}$  be a hypothesis space of VC dimension  $d$  and  $h : \mathcal{Z} \mapsto \mathcal{Y}, \forall h \in \mathcal{H}$ . When using DBE, given a feature extraction function  $\mathcal{F}^g : \mathcal{X} \mapsto \mathcal{Z}$  that shared between  $\mathcal{D}_i$  and  $\mathcal{D}$ , a random labeled sample of size  $m$  generated by applying  $\mathcal{F}^g$  to a random sample from  $\mathcal{U}_i$  labeled according to  $c^*$ , then for every  $h^g \in \mathcal{H}$ , with probability at least  $1 - \delta$ :

$$\mathcal{L}_{\mathcal{D}}(h^g) \leq \mathcal{L}_{\hat{\mathcal{D}}_i}(h^g) + \sqrt{\frac{4}{m} \left( d \log \frac{2em}{d} + \log \frac{4}{\delta} \right)} + d_{\mathcal{H}}(\tilde{\mathcal{U}}_i^g, \tilde{\mathcal{U}}^g) + \lambda_i,$$

where  $\mathcal{L}_{\hat{\mathcal{D}}_i}$  is the empirical loss on  $\mathcal{D}_i$ ,  $e$  is the base of the natural logarithm, and  $d_{\mathcal{H}}(\cdot, \cdot)$  is the  $\mathcal{H}$ -divergence between two distributions.  $\lambda_i := \min_{h^g} \mathcal{L}_{\mathcal{D}}(h^g) + \mathcal{L}_{\mathcal{D}_i}(h^g)$ ,  $\tilde{\mathcal{U}}_i^g \subseteq \mathcal{Z}$ ,  $\tilde{\mathcal{U}}^g \subseteq \mathcal{Z}$ , and  $d_{\mathcal{H}}(\tilde{\mathcal{U}}_i^g, \tilde{\mathcal{U}}^g) \leq d_{\mathcal{H}}(\tilde{\mathcal{U}}_i, \tilde{\mathcal{U}})$ .  $\tilde{\mathcal{U}}_i^g$  and  $\tilde{\mathcal{U}}^g$  are the induced distributions of  $\mathcal{U}_i$  and  $\mathcal{U}$  under  $\mathcal{F}^g$ , respectively.  $\tilde{\mathcal{U}}_i$  and  $\tilde{\mathcal{U}}$  are the induced distributions of  $\mathcal{U}_i$  and  $\mathcal{U}$  under  $\mathcal{F}$ , respectively.  $\mathcal{F}$  is the feature extraction function in the original FedAvg without DBE.

# [FL]: HtFLlib, heterogeneous FL algorithm library

- Beginner-friendly
- PFLlib compatible
- Extensible



# [FL]: HtFLlib on Device

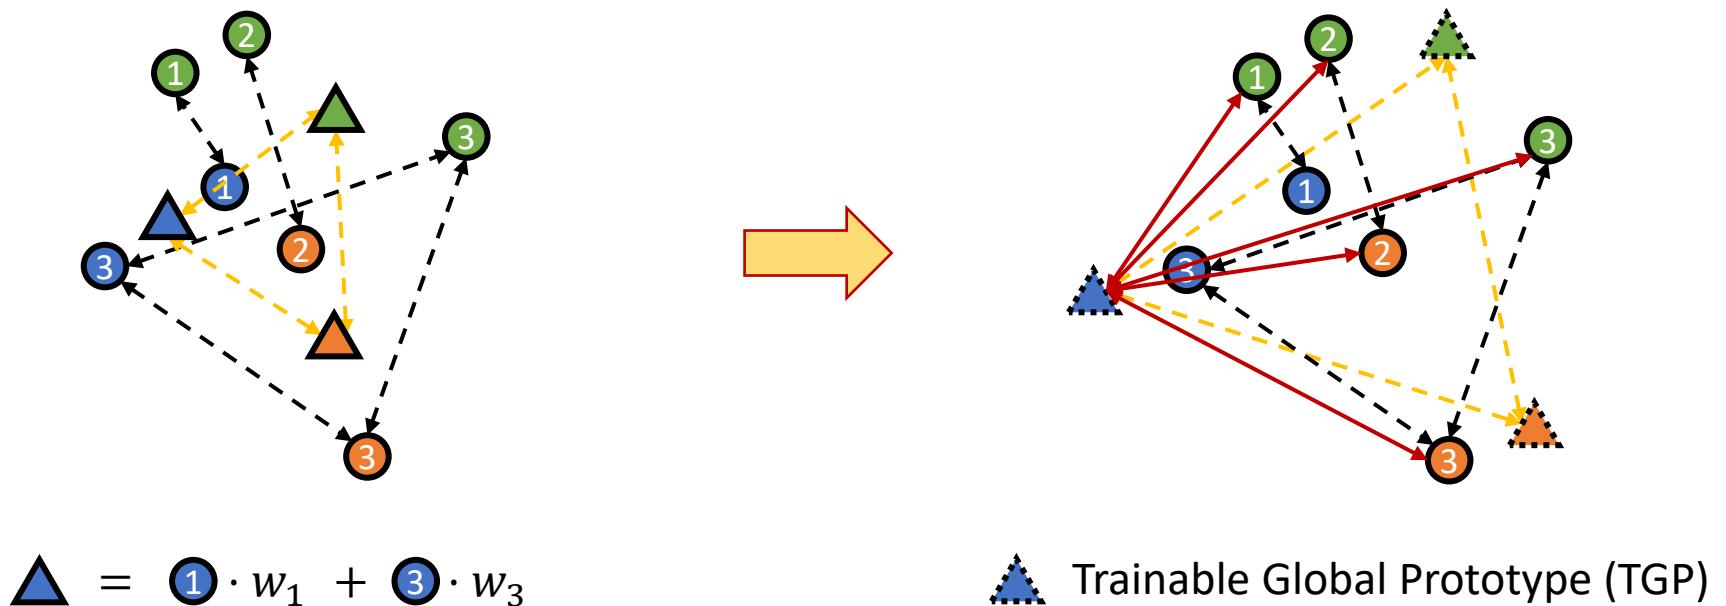
- Real-world deployment of HtFL methods
  - + CoLEXT, + real-world datasets, + systematical metrics



- 28 Single Board Computers (SBC)
  - Orange Pi, LattePanda, Nvidia Jetson
- 20 Smartphones
  - Samsung, Xiaomi, Google Pixel, Asus ROG, One Plus
- High Voltage Power Meter
- Wired and wireless networking
- Workstation - FL Server

# [FL]: FedTGP (Feature Align, Knowledge Distill)

- Remove weighted-averaging
- Consider the uploaded client prototypes as data
- **Enlarge** the global prototype margin



# [FL]: FedTGP (Feature Align, Knowledge Distill)

- Train global prototypes using **Adaptive-margin-enhanced Contrastive Learning (ACL)**
- ACL is **universal** and can be applied to other tasks

$$\min_{\hat{\mathcal{P}}} \sum_{c=1}^C \mathcal{L}_P^c,$$

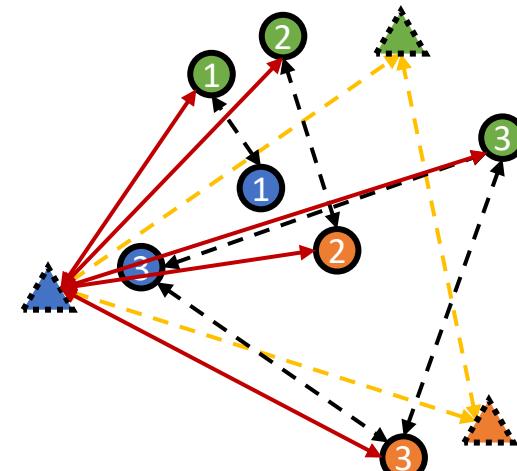
$$\mathcal{L}_P^c = \sum_{i \in \mathcal{I}^t} -\log \frac{e^{-(\phi(P_i^c, \hat{P}^c) + \delta(t))}}{e^{-(\phi(P_i^c, \hat{P}^c) + \delta(t))} + \sum_{c'} e^{-\phi(P_i^c, \hat{P}^{c'})}}$$

$$\delta(t) = \min(\max_{c \in [C], c' \in [C], c \neq c'} \phi(Q_t^c, Q_t^{c'}), \tau),$$

$$Q_t^c = \frac{1}{|\mathcal{P}_t^c|} \sum_{i \in \mathcal{I}^t} P_i^c, \forall c \in [C]$$

$\tau$  is a margin threshold

maximum cluster margin



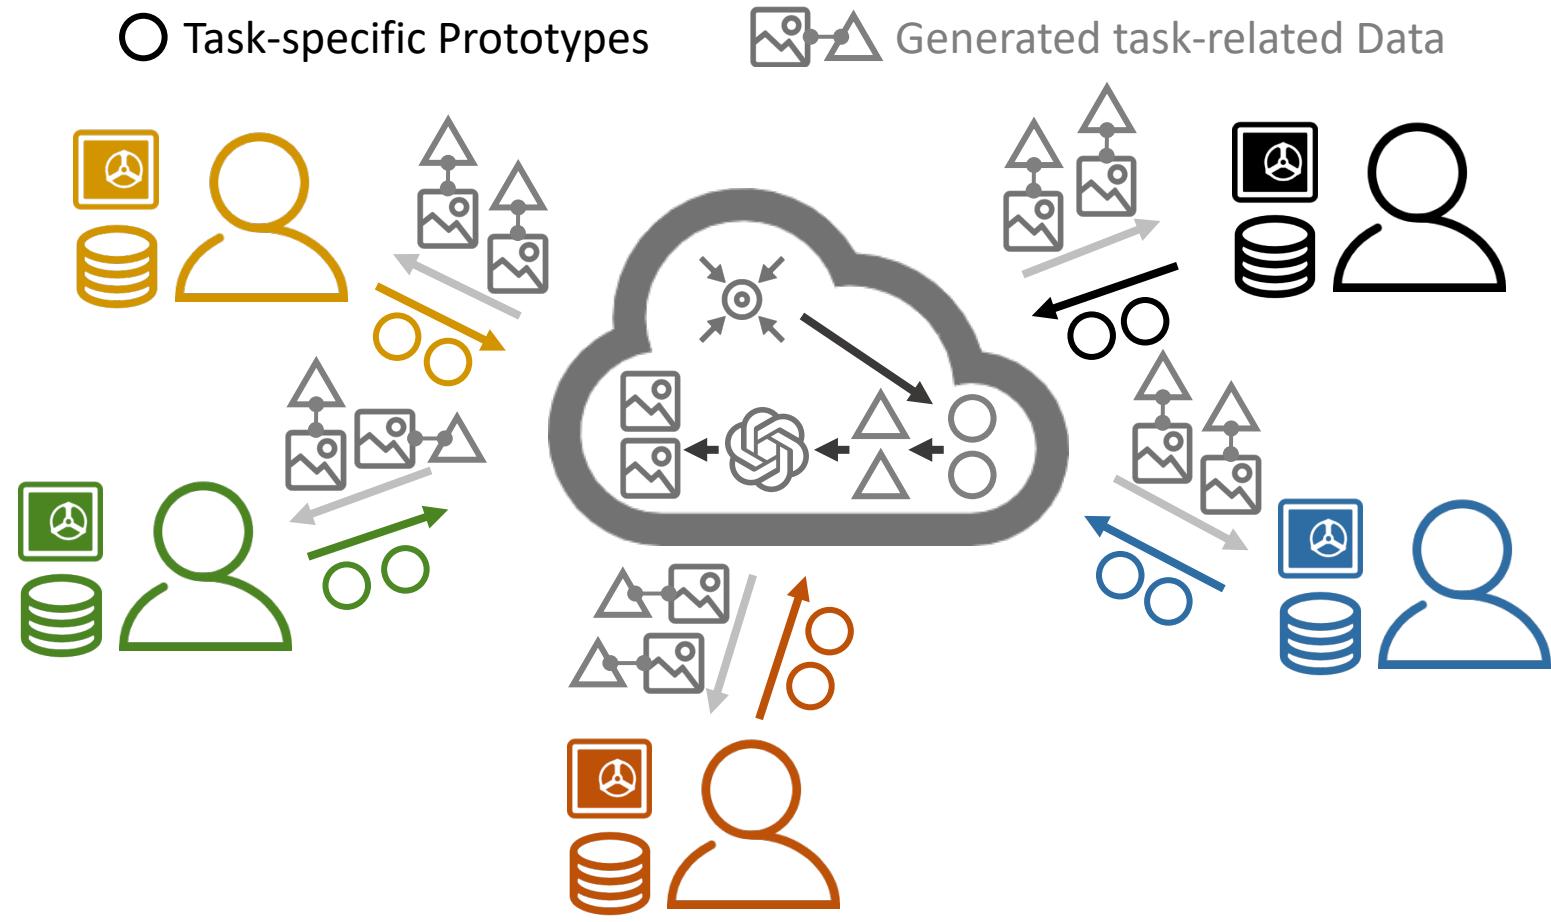
$\triangle \hat{P}^c$ : A TGP of class  $c$

$\hat{\mathcal{P}}$ : All TGP

$\bullet P_i^c$ : A prototype of class  $c$  from client  $i$

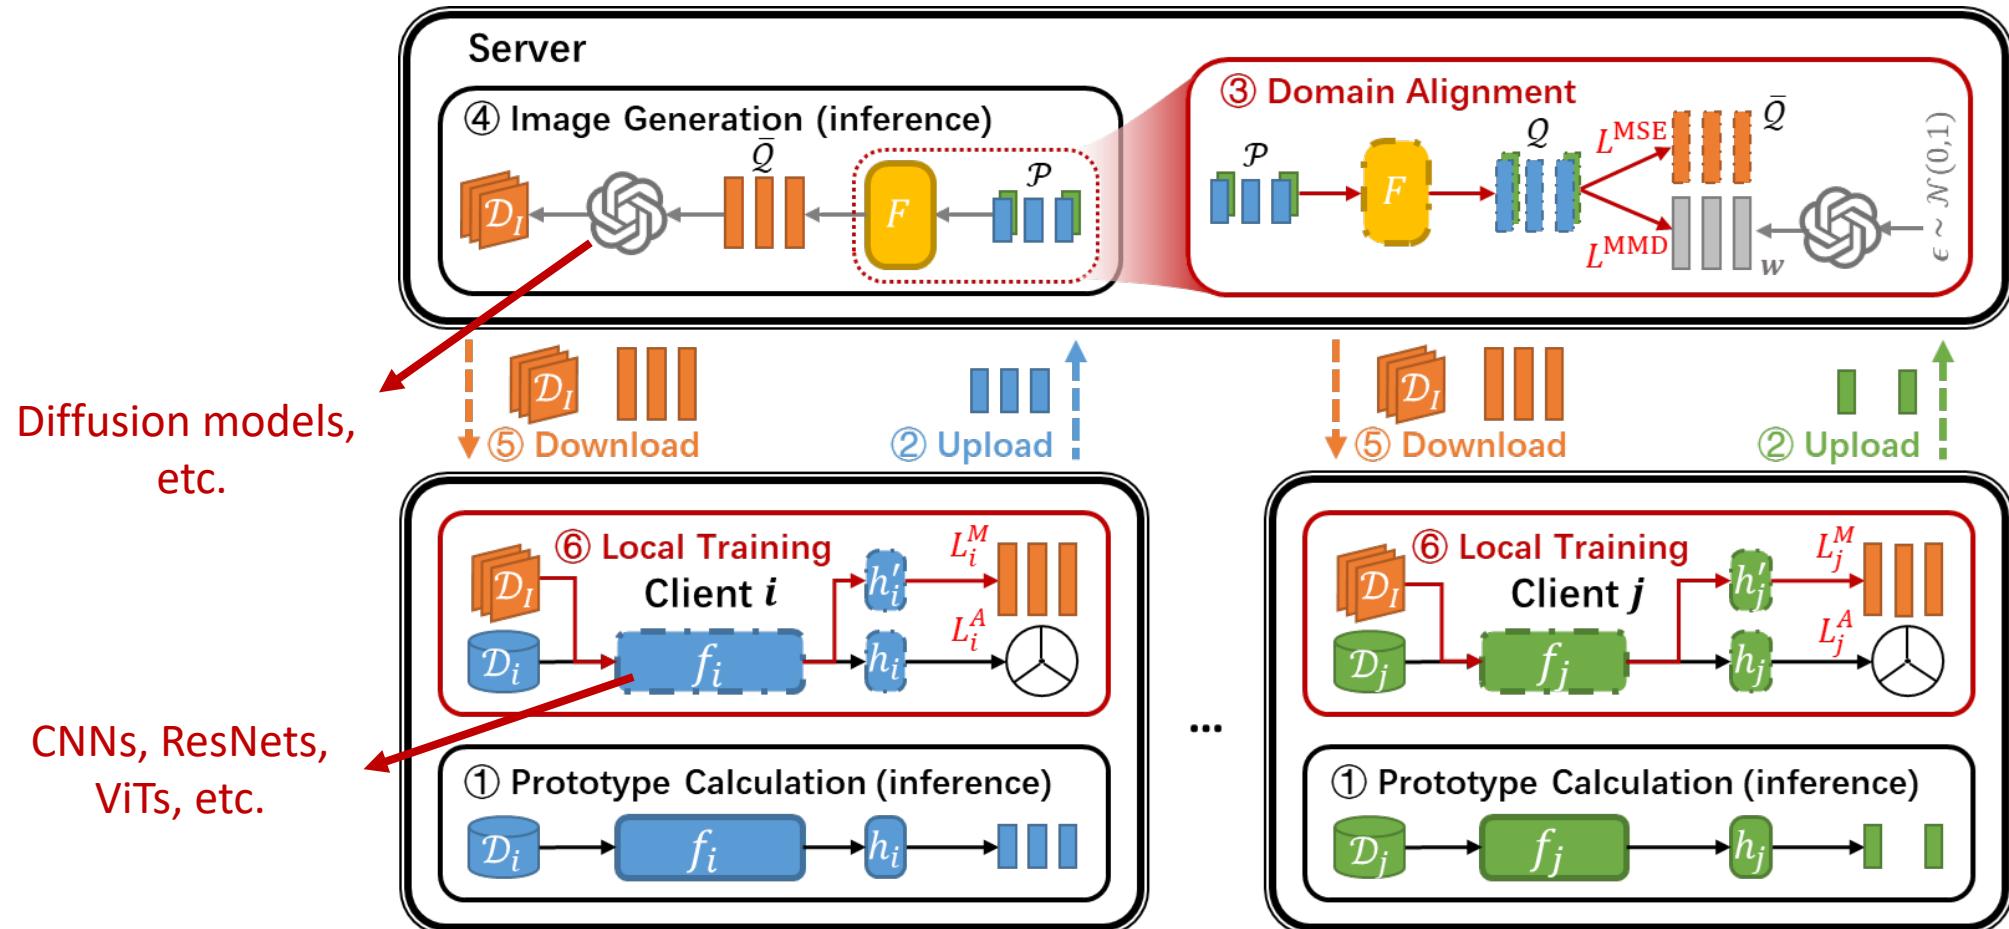
# [FL]: FedKTL (Feature Align, Knowledge Distill)

- Transfer **common knowledge** from the generator to clients
- Obtain **task-specific knowledge** from other clients



# [FL]: FedKTL (Feature Align, Knowledge Distill)

- Align small models' feature space with the generative model's
- Transfer global knowledge using an **additional supervised local task**



# [FL]: FedKTL (Feature Align, Knowledge Distill)

- FedKTL can **adapt to various generators** that were pre-trained using various datasets
- The **semantics of the generated images** can be different from clients' data



(a) Client #1



(b) AFHQv2



(c) Benches



(d) FFHQ-U



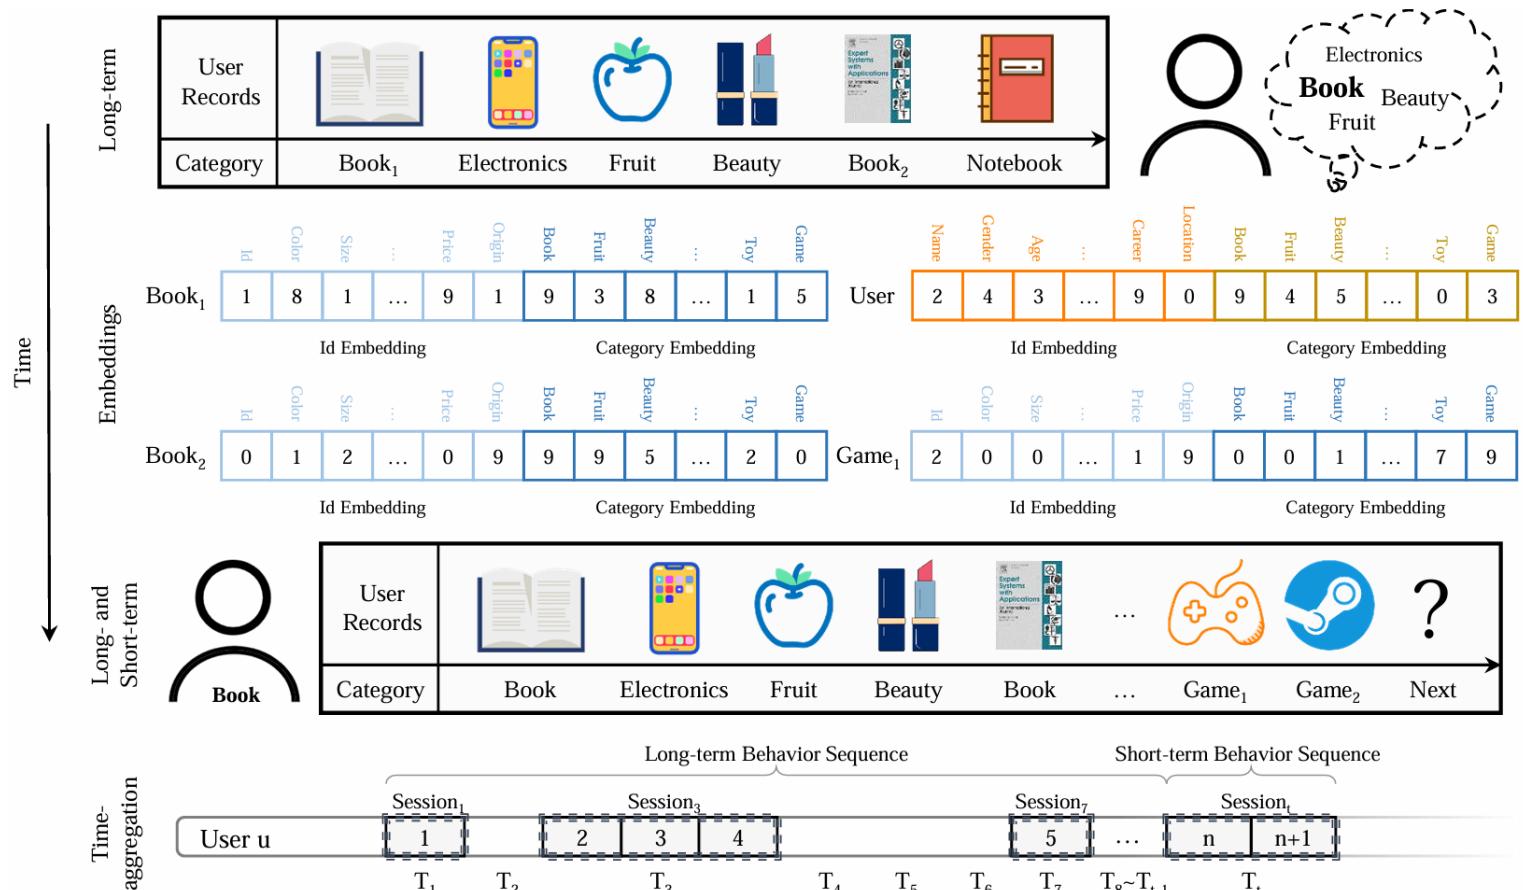
(e) WikiArt

Generators pre-trained on different image datasets

	$\lambda = 0.05$	$\lambda = 0.1$	$\lambda = 0.5$
AFHQv2	$26.82 \pm 0.32$	<b><math>27.05 \pm 0.26</math></b>	$26.32 \pm 0.52$
Bench	$27.71 \pm 0.25$	<b><math>28.36 \pm 0.42</math></b>	$27.56 \pm 0.50$
FFHQ-U	<b><math>27.28 \pm 0.23</math></b>	$27.21 \pm 0.35$	$26.59 \pm 0.47$
WikiArt	$27.37 \pm 0.51$	<b><math>27.48 \pm 0.33</math></b>	$27.30 \pm 0.15$

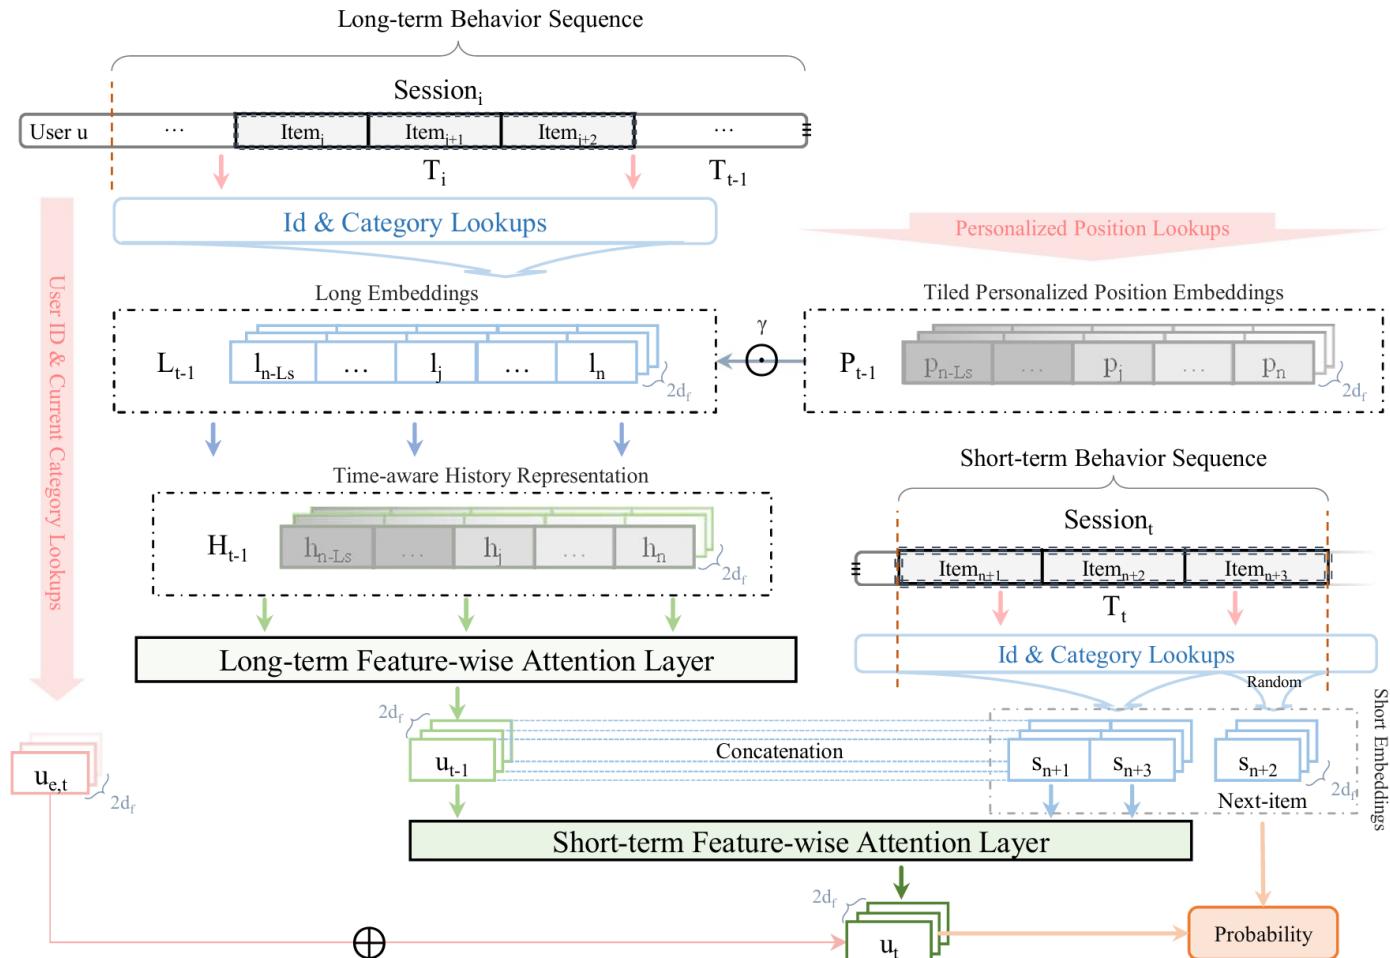
# Recommender System (RS)

- TLSAN: Time-aware Long- and Short-term Attention Network for Next-item Recommendation
- Users have **personalized taste for time**



# [RS]: TLSAN (Personalization)

- Capture personalized time-aggregation pattern in long-term attention



# Feel free to contact me!

Home page: <https://github.com/TsingZ0>



Thanks!